

Introduction to Optimization

Lecture 04: Strict and strong convexity. Iterative algorithms. Descent methods.



Characterizations of differentiable convex functions

Proposition

Let $f : D \subset \mathbb{R}^N \rightarrow \mathbb{R}$ be differentiable. The following are equivalent:

- ① f is convex;
- ② for all $x, y \in D$, $f(y) \geq f(x) + \nabla f(x) \cdot (y - x)$;
- ③ for all $x, y \in D$, $(\nabla f(y) - \nabla f(x)) \cdot (y - x) \geq 0$.

If f is twice differentiable, the three statements above are equivalent to

- ④ for all $x \in D$, $\nabla^2 f(x)$ is positive semidefinite.

Strict and strong convexity

A function $f : D \subset \mathbb{R}^N \rightarrow \mathbb{R}$ is **strictly convex** if D is convex and

$$f(\lambda x + (1 - \lambda)y) < \lambda f(x) + (1 - \lambda)f(y)$$

for all $x, y \in D$ and all $\lambda \in (0, 1)$

Strict and strong convexity

A function $f : D \subset \mathbb{R}^N \rightarrow \mathbb{R}$ is **strictly convex** if D is convex and

$$f(\lambda x + (1 - \lambda)y) < \lambda f(x) + (1 - \lambda)f(y)$$

for all $x, y \in D$ and all $\lambda \in (0, 1)$, and it is **strongly convex** if D is convex and

$$f(\lambda x + (1 - \lambda)y) \leq \lambda f(x) + (1 - \lambda)f(y) - \frac{\alpha}{2}\lambda(1 - \lambda)\|x - y\|^2$$

for all $x, y \in D$ and all $\lambda \in (0, 1)$.




Exercises

- 1 Find examples of: a function that is convex, but not strictly convex; and a function that is strictly convex, but not strongly convex.
- 2 When is the function $f(x) = \frac{1}{2}\|Ax - y\|^2$ strictly/strongly convex?
- 3 Prove that every strictly convex function $f : \mathbb{R}^N \rightarrow \mathbb{R}$ has at most one minimizer, and every strongly convex function $f : \mathbb{R}^N \rightarrow \mathbb{R}$ has exactly one minimizer.
- 4 Can you obtain characterizations of strict and strong convexity of f in terms of properties of ∇f and $\nabla^2 f$?


Introduction to iterative algorithms

Example: Bisection method to solve $g(x) = 0$




Introduction to iterative algorithms

Example: Bisection method to solve $g(x) = 0$




Introduction to iterative algorithms

Example: Bisection method to solve $g(x) = 0$




Introduction to iterative algorithms

Example: Bisection method to solve $g(x) = 0$




Introduction to iterative algorithms

Example: Bisection method to solve $g(x) = 0$




Introduction to iterative algorithms

Example: Bisection method to solve $g(x) = 0$




Introduction to iterative algorithms

Example: Bisection method to solve $g(x) = 0$




Introduction to iterative algorithms

Example: Bisection method to solve $g(x) = 0$




Introduction to iterative algorithms

Example: Bisection method to solve $g(x) = 0$



Introduction to iterative algorithms

Example: Bisection method to solve $g(x) = 0$



After k iterations, the distance to a solution is $|x_k - \hat{x}| \leq 2^{-k}$.

Introduction to iterative algorithms

An **iterative algorithm** is a procedure that computes a sequence (x_n) of points in \mathbb{R}^N that approximate a solution to a problem. It requires:

- An initial guess x_0 .
- A sequence (p_k) of parameters (typically $p_k \in \mathbb{R}^M$ for all $k \geq 0$).
- An operator $T : \mathbb{R}^N \times \mathbb{R}^M \rightarrow \mathbb{R}^N$ used to compute x_{k+1} , given x_k :

$$x_{k+1} = T(p_k, x_k).$$

- A stopping rule that is activated when the approximation is **sufficiently good**.

Introduction to iterative algorithms

An **iterative algorithm** is a procedure that computes a sequence (x_n) of points in \mathbb{R}^N that approximate a solution to a problem. It requires:

- An initial guess x_0 .
- A sequence (p_k) of parameters (typically $p_k \in \mathbb{R}^M$ for all $k \geq 0$).
- An operator $T : \mathbb{R}^N \times \mathbb{R}^M \rightarrow \mathbb{R}^N$ used to compute x_{k+1} , given x_k :

$$x_{k+1} = T(p_k, x_k).$$

- A stopping rule that is activated when the approximation is **sufficiently good**.

Important questions: **convergence** and **complexity**.

Stopping rules in minimization problems

Ideally, the algorithm should stop when this is when

- x_k is close to a minimizer, ✗
- $f(x_k)$ is close to the optimal value $\inf(f)$, ✗
- $\|\nabla f(x_k)\|$ is small. ✓

Stopping rules in minimization problems

Ideally, the algorithm should stop when this is when

- x_k is close to a minimizer, ✗
- $f(x_k)$ is close to the optimal value $\inf(f)$, ✗
- $\|\nabla f(x_k)\|$ is small. ✓

Common implementable rules:

- $\|\nabla f(x_k)\| < \varepsilon$
- $f(x_{k+1}) - f(x_k) < \varepsilon$
- $\|x_{k+1} - x_k\| < \varepsilon$

Stopping rules in minimization problems

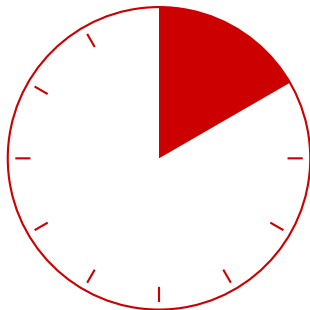
Ideally, the algorithm should stop when this is when

- x_k is close to a minimizer, ✗
- $f(x_k)$ is close to the optimal value $\inf(f)$, ✗
- $\|\nabla f(x_k)\|$ is small. ✓

Common implementable rules:

- | | | |
|---|---|---|
| • $\ \nabla f(x_k)\ < \varepsilon$ | • $f(x_{k+1}) - f(x_k) < \varepsilon$ | • $\ x_{k+1} - x_k\ < \varepsilon$ |
| • $\frac{\ \nabla f(x_k)\ }{\ \nabla f(x_0)\ } < \varepsilon$ | • $\frac{f(x_{k+1}) - f(x_k)}{f(x_1) - f(x_0)} < \varepsilon$ | • $\frac{\ x_{k+1} - x_k\ }{\ x_1 - x_0\ } < \varepsilon$ |

Break



Descent methods

Many algorithms are based on the idea of (sufficient) descent: given x_k , find x_{k+1} such that

$$(1) \quad f(x_{k+1}) \leq f(x_k) - \delta_k^2.$$

Descent methods

Many algorithms are based on the idea of (sufficient) descent: given x_k , find x_{k+1} such that

$$(1) \quad f(x_{k+1}) \leq f(x_k) - \delta_k^2.$$

One way is to find $d_k \in \mathbb{R}^N$ and $\alpha_k > 0$, such that (1) holds with

$$x_{k+1} = x_k - \alpha_k d_k.$$

Descent methods

Many algorithms are based on the idea of (sufficient) descent: given x_k , find x_{k+1} such that

$$(1) \quad f(x_{k+1}) \leq f(x_k) - \delta_k^2.$$

One way is to find $d_k \in \mathbb{R}^N$ and $\alpha_k > 0$, such that (1) holds with

$$x_{k+1} = x_k - \alpha_k d_k.$$

We say $-d_k$ is a **descent direction**, and α_k is the **step size**, **step length** or **learning rate** (in ML).

Motivation: 3 case studies

Let us analyze the behavior in the following cases:

❶ $f(x) = x^2$

Motivation: 3 case studies

Let us analyze the behavior in the following cases:

① $f(x) = x^2$

② $f(x) = 1/x, \text{ dom}(f) = (0, \infty).$

Motivation: 3 case studies

Let us analyze the behavior in the following cases:

① $f(x) = x^2$

② $f(x) = 1/x, \text{ dom}(f) = (0, \infty).$

③ $f(x) = 1/x, \text{ dom}(f) = (-\infty, 0).$

L -smoothness

If there is time

A differentiable function $f : A \subset \mathbb{R}^N \rightarrow \mathbb{R}$ is **L -smooth**, with $L > 0$, if

$$\|\nabla f(x) - \nabla f(y)\| \leq L\|x - y\|$$

for all $x, y \in A$.

L -smoothness

If there is time

A differentiable function $f : A \subset \mathbb{R}^N \rightarrow \mathbb{R}$ is **L -smooth**, with $L > 0$, if

$$\|\nabla f(x) - \nabla f(y)\| \leq L\|x - y\|$$

for all $x, y \in A$.

Proposition (Descent Lemma)

If f is L -smooth and A is convex, then

$$|f(y) - f(x) - \nabla f(x) \cdot (y - x)| \leq \frac{L}{2}\|x - y\|^2$$

for all $x, y \in A$.

