

Information Security

(WBCS004-05)

Fatih Turkmen

Some slides are borrowed from Dr. Frank B. Brokken

Today

- Hashing:
 - Concept and requirements
 - Collisions
 - Use of hashing in cryptography
 - CRC: a non-cryptographic hash
 - MD5, SHA-x, Tiger
 - Sharing Secrets / Key Escrow
 - Information Hiding
 - E-mail peculiarities

Hashing



*Yes, the alchemists worshipped the antimatter, vowing dark,
agile actions while announcing algebra...*

Hashing

- Why do we hash?
 - Verification of the integrity of a message
 - Authentication
 - Message fingerprinting
 - Digital Signatures

Requirements

- *Compression*: A hashing function h computes a small string/number matching a large piece of information (e.g., a message).
- *Efficiency*: Computing $h(x)$ must be easy/fast.
- *Trap-door*: Given $h(x)$, x cannot be “easily” retrieved
- *Collision resistant*: Infeasible to find y for which $h(x) == h(y)$ (x is given or freely selectable)
 - *Collisions do exist*: If $h()$ results in N bits and if x consists of M bits ($M > N$) then there must exist 2^{M-N} collisions How many collisions?

Hashing Algorithm

Try 1: Simple addition: If x_i are bytes, $h(x) = \sum x_i \pmod{256}$

✓ Compresses, easy to compute, cannot be inverted.

✗ Unfortunately: many **collisions**

Allowing blanks, here are some for hashes

icy porn net → (69 63 79 20 70 6f 72 6e 20 6e 65 74) = X + 20 (mod 256)

inept crony → (69 6e 65 70 74 20 63 72 6f 6e 79) = X (mod 256)

intern copy → (69 6e 74 65 72 6e 20 63 6f 70 79) = X (mod 256)

no inept cry → (6e 6f 20 69 6e 65 70 74 20 63 72 79) = X + 20 (mod 256)

ASCII ¹



¹ <https://commons.wikimedia.org/wiki/File:ASCII-Table-wide.svg>

Hashing

Try 2: Modification: If x_i are bytes, multiply the values w
index $h(x) = \sum i * x_i \pmod{256}$

✓ Compresses, easy to compute, cannot be inverted.

✓ Fewer collisions (i.e., better distribution over the "has

✗ it's still easy to construct collisions:

	1	*	a	+	2	*	b	=	h()
@0	1	*	64	+	2	*	48	=	160
>1	1	*	62	+	2	*	49	=	160



Given this scheme, what are the hashes of "@0" and ">1"?

ASCII printable characters					
32	space	64	@	96	`
33	!	65	A	97	a
34	"	66	B	98	b
35	#	67	C	99	c
36	\$	68	D	100	d
37	%	69	E	101	e
38	&	70	F	102	f
39	'	71	G	103	g
40	(72	H	104	h
41)	73	I	105	i
42	*	74	J	106	j
43	+	75	K	107	k
44	,	76	L	108	l
45	-	77	M	109	m
46	.	78	N	110	n
47	/	79	O	111	o
48	0	80	P	112	p
49	1	81	Q	113	q
50	2	82	R	114	r
51	3	83	S	115	s
52	4	84	T	116	t
53	5	85	U	117	u
54	6	86	V	118	v
55	7	87	W	119	w
56	8	88	X	120	x
57	9	89	Y	121	y
58	:	90	Z	122	z
59	;	91	[123	{
60	<	92	\	124	
61	=	93]	125	}
62	>	94	^	126	~
63	?	95	_		

Collisions

Read about the **Birthday Paradox** in the book!

Tiger Hash

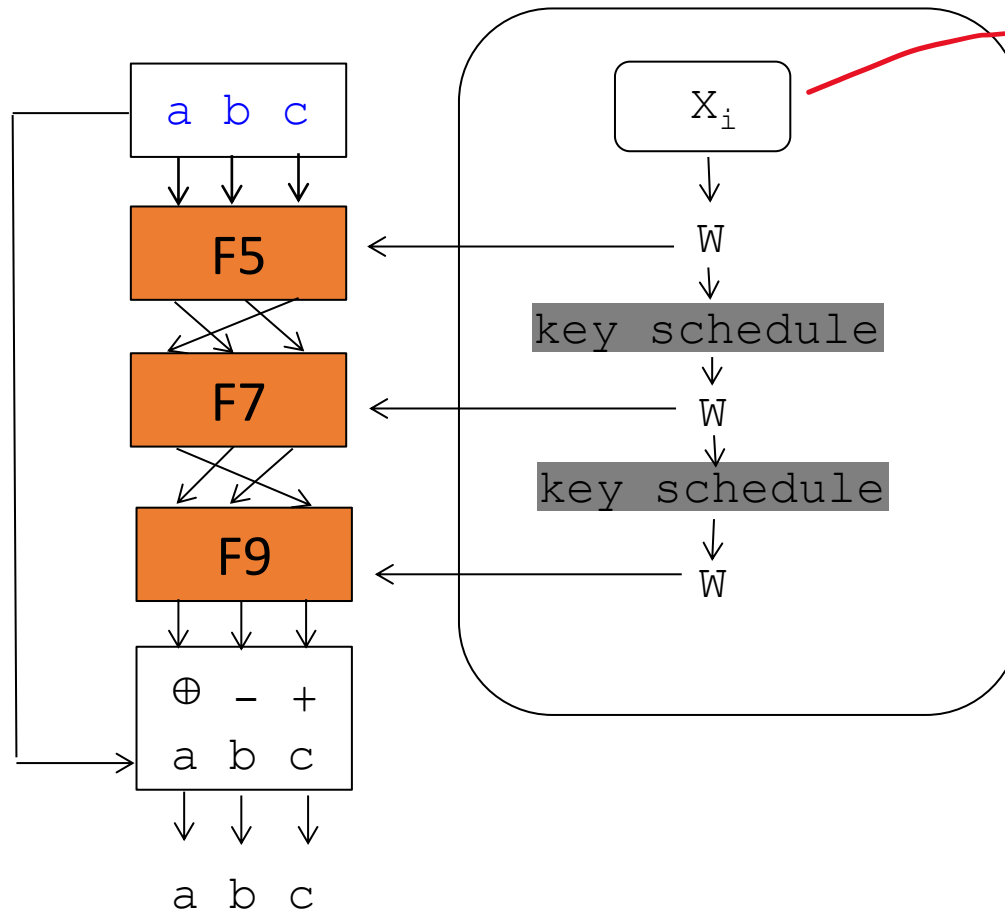
- Developed (1995) by Ross Anderson & Eli Biham.
- Resembles block ciphers
- Operates on blocks of 512 bits (padding may be applied if necessary)
- Resulting hash value (digest): 192 bits (works well with 64-bit processors)
- 4 S-boxes mapping 8 bits to 64 bits
- Uses a *key schedule*, using the input blocks as key.
- Tiger applies one *outer round* on each 512-bit block.



Tiger Hash (Outer Rounds)



- Input is $X = (X_0, X_1, \dots, X_{n-1})$
- Tiger's **outer round**: Applied to each 512-bit block (i.e. X_i):




Note: The message itself is used as a key, since there is no key!

- There are n iterations of the outer round
- Initial a, b, c have fixed values (e.g., a is 0x123456789abcdef)

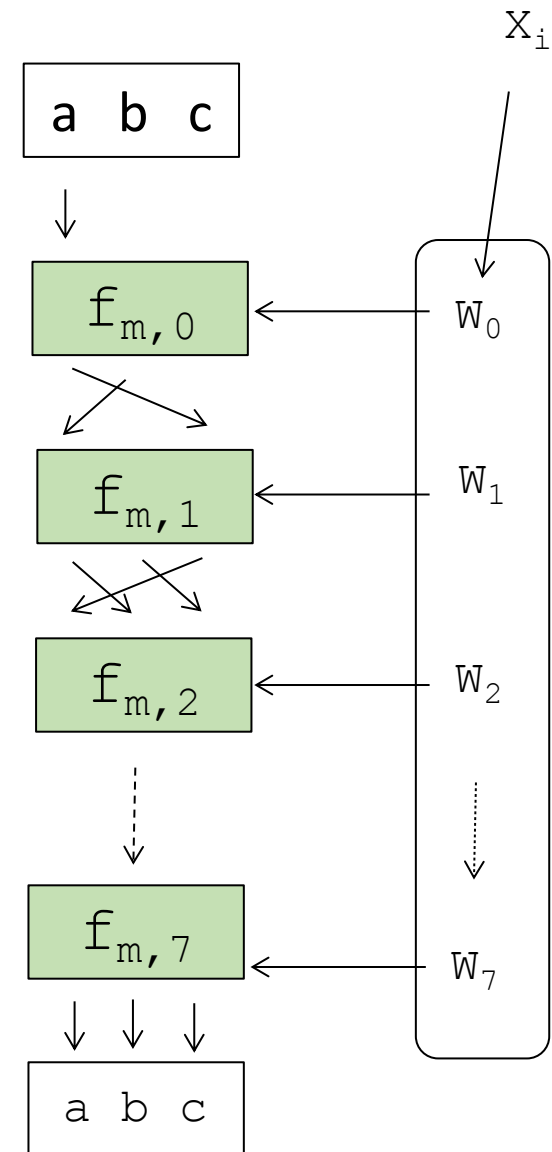
Tiger Hash (Outer Rounds)

- Three outer round functions F_5 , F_7 , F_9
- a , b , c : each 64 bits
- a leaving F_5 becomes b of F_7 , b leaving F_7 becomes a of F_9 etc...
- *Final* a , b , c is the hash value, thus the final output is 192 bits



Tiger Hash (Inner Rounds)

- Each F_m consists of 8 *inner rounds* where $m \in \{5, 7, 9\}$:
 - Each w_i is a 64-bit section of a 512 bit input block, i.e. $W = (w_0, w_1, \dots, w_{n-1})$
 - Each $f_{m,i}$ receives a permutation of the a,b,c output by $f_{m,i-1}$.
E.g., $(abc), (bca), (cab)$:
 - $f_{m,0}$ to $f_{m,1}$: output b becomes input a



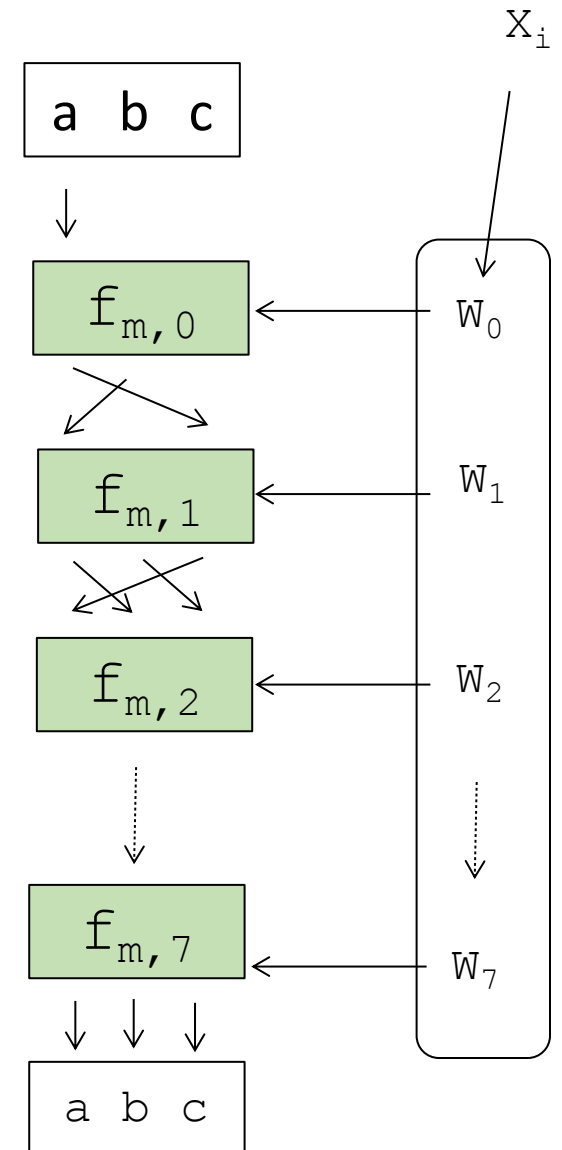
Tiger Hash (Inner Rounds)

- Final step in inner rounds
 - The 64 c-bits are split in 8 bits ($c_0 \dots c_7$):

$$\begin{array}{l}
 c \oplus = w_i \\
 a -= S[0][c_0] \oplus S[1][c_2] \oplus S[2][c_4] \oplus S[3][c_6] \\
 b += S[3][c_1] \oplus S[2][c_3] \oplus S[1][c_5] \oplus S[0][c_7] \\
 b *= m
 \end{array}$$

- The *key schedule* recomputes w_0 to w_7 between the f_m boxes (how? Next slide!)
 - cf. Stamp, p. 132.


An S-box element



Tiger Hash

- Tiger's key schedule, simplifying Table 5.1 (see Stamp, p. 93 for more details):


```
w0 -= w7  $\oplus$  0xa5a5a5a5a5a5a5a5;  
w1  $\oplus$ = w0;  
w2 += w1;  
w3 -= w2  $\oplus$  (~w1 << 19);  
w4  $\oplus$ = w3;  
w5 += w4;  
w6 -= w5  $\oplus$  (~w4 >> 23);  
w7  $\oplus$ = w6;  
w0 += w7;  
w1 -= w0  $\oplus$  (~w7 << 19);  
w2  $\oplus$ = w1;  
w3 += w2;  
w4 -= w3  $\oplus$  (~w2 >> 23);  
w5  $\oplus$ = w4;  
w6 += w5;  
w7 -= w6  $\oplus$  0x0123456789abcdef;
```



Tiger Hash

- The final step is called feedforward
- The results, say a' , b' , c' of F9 are XORed, subtracted and added with the initial a , b , c respectively.

- For the original proposal see Tiger Hash paper:
https://link.springer.com/content/pdf/10.1007/3-540-60865-6_46.pdf
- For the Cryptanalysis see:
<https://iacr.org/archive/asiacrypt2007/48330539/48330539.pdf>



Refresher on Message Integrity

- Use of cryptography for “Unauthorized Modification” (not about unauthorized reading!) of the plain text
- Message Authentication Code (MAC) (chapter 3.4)



What “principal” is this about?

- How does MAC work?

- Symmetric encryption, i.e., the same encryption key is used
- It works in **CBC mode**, i.e., blocks of messages M_0, M_1, \dots, M_{n-1}

$$C_0 = E(M_0 \oplus IV, K), \quad C_1 = E(M_1 \oplus C_0, K), \quad \dots, \quad C_{n-1} = E(M_{n-1} \oplus C_{n-2}, K)$$

- C_{n-1} , also called as CBC residue, serves as the MAC. The rest is discarded for the case of “integrity”.

(Example) Uses of Hashing: Integrity

- Verification of the **integrity** of a message
 - With MAC $\{c_0, c_1, \dots, c_{n-1}\} + c_{n-1}$ is sent for confidentiality + integrity.
- Requirements:
 - M and $h(M)$: changing M changes $h(M)$ and v.v.
 - M and $h(M)$ will be sent together since we are interested in integrity...

HMAC

- Integrity with respect to original message “M” must be protected (i.e., the figure).
- Enter: **Hashed Message Authentication Code**.
 - Prevent the change of hash!
 - Hashing functions typically process blocks of bytes.
 - We could prepend (or append) a key K to the message M (i.e., start with the key, or start with the message block(s))

Alice




$h(M), M$



Bob



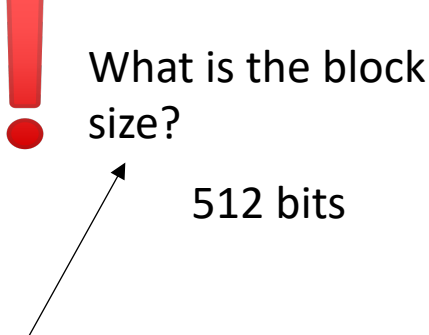
$h(M'), M'$



Trudy



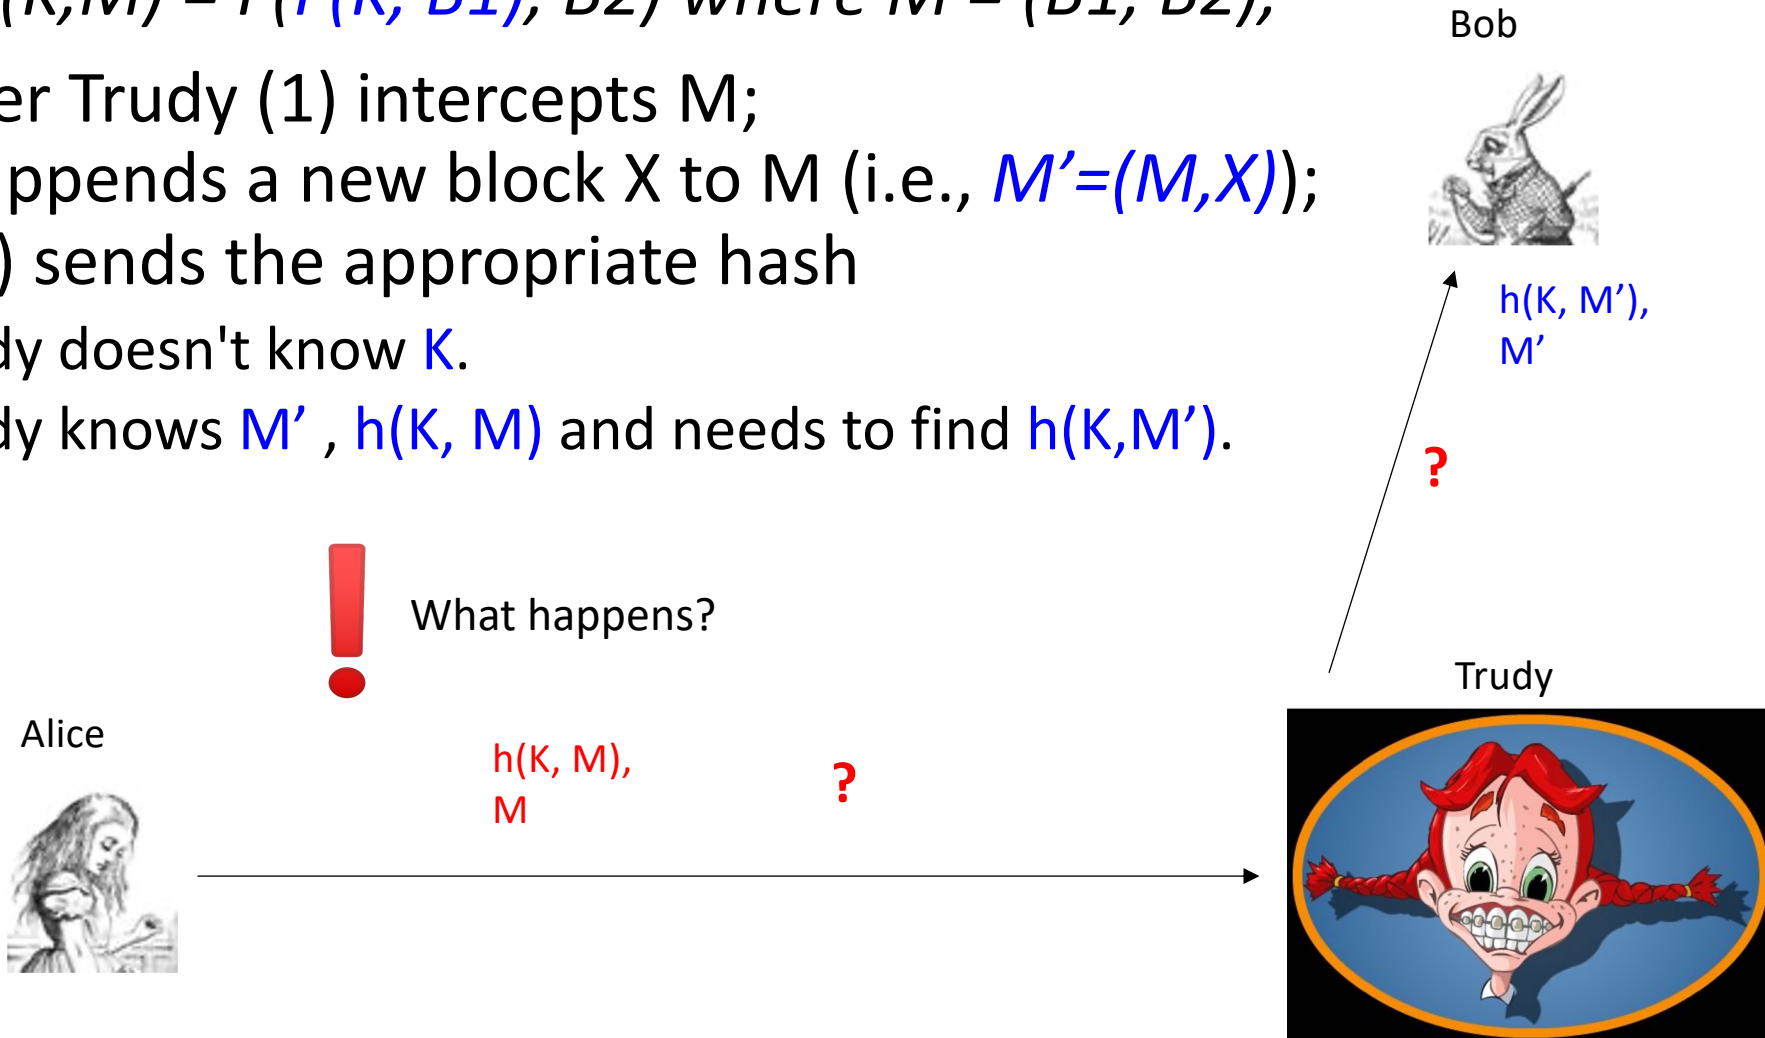
HMAC (cont.)

- So, should we use $h(K, M)$ or $h(M, K)$?
 - Let's consider the prepend: $h(K, M)$
 - Hash functions F tend to use *blocks* (e.g., *Tiger uses ?*), e.g., $M = (B1, B2)$.
 - The common case is to use the hash of the previous block as input when computing the next block's hash:
 - $h(M) = F(F(Init, B1), B2)$, where F is similar to the outer round of Tiger
- 
- What is the block size?
512 bits

HMAC (cont.)

Given $h(K, M) = F(F(K, B1), B2)$ where $M = (B1, B2)$,

- Intruder Trudy (1) intercepts M ;
 (2) appends a new block X to M (i.e., $M' = (M, X)$);
 and (3) sends the appropriate hash
 - Trudy doesn't know K .
 - Trudy knows M' , $h(K, M)$ and needs to find $h(K, M')$.



HMAC (cont.)

- Use $h(K, M)$? Bad idea...
- $h(K, M) = F(F(K, B1), B2)$
 - If Trudy appends X and sends M' not knowing K:

- $h(K, M, X) =$

$$F(F(F(A, K), M), X) = F(h(K, M), X)$$

Intercepted
by Trudy earlier

Set A=0 if you use Tiger
which does not have key
but fixed constants..

Bob



$h(K, M'),$
 M'

Trudy



HMAC (cont.)

$$h(M1, K) = h(h(M1), K) = h(h(M2), K)$$

- Use $h(M, K)$ instead.
- Less serious but (if there is) a known *collision* ($h(M1) == h(M2)$) renders the hash function *insecure*.
 - Note: M1 and M2 need to be a multiple of the block size.
 - This happened to MD5. SHA1 by now is also considered insecure.
 - Better to use HMAC described in RFC 2104:
(<https://www.ietf.org/rfc/rfc2104.txt>)
 - K, unknown to Trudy, is required to finalize the hash computation

Look for a nice discussion here! :

<https://stackoverflow.com/questions/7885268/simple-enquiry-on-hash-algorithm>

HMAC (cont.)

- $h(M,K)$ is preferred over $h(K,M)$
- None of these solutions is complete safe!
 - RFC 2104 offers a solution, B the block length (i.e., $512/8=64$)
 - Thoroughly mixing the key to the hash!

B: hash block size **in bytes** (e.g., $B = 64$)
define: *ipad* = 0x36 repeated B times
opad = 0x5C bytes repeated B times.

$$\mathbf{HMAC(M,K) = h_1(K \oplus opad, h_2(K \oplus ipad, M))}$$

- (ipad and opad could be omitted)
- Note: h_1 and h_2 are the same hash function. h_2 is the real work but it reduces the message to digest/hash so h_1 is quickly computed thereafter.
- See also <https://en.wikipedia.org/wiki/HMAC>

Example Non-cryptographic Uses of Hashing: CRC

- Cyclic Redundancy Check (CRC)
 - Not a cryptographically-acceptable hash function.
 - Intended for networking applications: detecting transmission errors.
 - WEP uses (inappropriately) CRCs.

Computation of CRC

Given a divisor of size n ;

- Add $n-1$ 0-bits at the end of the dividend (i.e., data stream).
 - Once the first (leftmost) bit of the input stream is 1: *xor* the leftmost n bits of the input stream by the divisor:
 - Continue the process until the remainder is 0 or smaller than the divisor
- The *remainder* ($n-1$ bits) is the CRC.

```
bitstream: 101010110000
divisor:   10011
           11001
```

Computation of CRC

- The remainder (*n-1* bits) is the CRC:

$$\boxed{x \mid 0} \oplus d = \boxed{0 \mid c} \quad (\oplus d: d \text{ 'slides' over } x \mid 0)$$

- So:

$$\boxed{x \mid 0 \oplus d = c}$$

- Consequently: the CRC of a bitstream + its CRC equals 0:

$$\boxed{x \mid c \oplus d = 0 \mid 0}$$

Computation of CRC

- Example:




The message: 10011100

The divisor: 10011

What is the CRC?

1100



See this for an explanation of the division (in CRC):
<https://www.youtube.com/watch?v=kscjEvjTVBI>

Collisions in CRC

- It's easy to create CRC collisions
 - Look at the final (intermediate) value before the CRC bits are added at the end of the bit stream:

bitstream: 101010111010

```

...
10101
10011
11010
    
```

bitstream:	101010110000	101010111010
divisor:	10011	10011
	11001	11001
	10011	10011
	10101	10101
	10011	10011
	11000	11010
	10011	10011
	10110	10011
	10011	10011
	1010	0

- Observation: Once 110 is the remainder of the division, the resulting CRC is 1010
- Earlier bits are irrelevant, as long as the result, ignoring the remainder, equals 110.

Collisions in CRC

- Finding CRC collisions
 - Change the bitpattern *ad lib*, and turn the final #divisor bits into .- characters, then solve for the dots. Originally:

```
bitstream: 101010111010
            (... )
            10101
            10011
            11010
```

Collisions in CRC

- Find a collision: $CRC = 1010$, $d = 10011$

original: 10101011

modified: 010.....1010

divisor: 10011
 0.....
 10011
 110

find 5 bits

bitstream: 101010110000
 divisor: 10011
 11001
 10011
 10101
 10011
 11000
 10011
 10110
 10011
 1010

101010111010
 10011
 11001
 10011
 10101
 10011
 11010
 10011
 10011
 10011
 0

Hashing

- Find a collision: $CRC = 10_d, d = 10011$

original: 10101011

modified: 010.....1010 010110011010

divisor: 10011 10011
 0..... 010101
 10011 10011
 00110 110


- the required bits are all implied and easy to find.

Hashing


- MD5, SHA-1
 - Both MD5 and SHA-1 were extremely popular
 - MDx (128 bit hash) hashes are now considered *insecure*, as collisions can be found.
 - SHA-1 (180 bit) is an improvement, but is in fact by now superseded by SHA-256.
(cf. <http://csrc.nist.gov/groups/ST/hash/policy.html>),
Software computing SHA-x is widely available

Hashing

- How insecure is SHA-1?
- *Schneier* reports that in approx. 2^{74} computer cycles a SHA-1 collision is found



Hashing




- How insecure is SHA-1?
- *Schneier* reports that in approx. 2^{74} computer cycles a SHA-1 collision is found
- In 2016 a *core* ran at approx. 2^{33} cycles/sec. Assume a *processor* has 8 cores, and a multi-processor *server* 4 processors; then a server did $2^{33+3+2} = 2^{38}$ cycles/second.

Hashing



- How insecure is SHA-1?
 - *Schneier* reports that in approx. 2^{74} computer cycles a SHA-1 collision is found
 - In 2016 a *core* ran at approx. 2^{33} cycles/sec. Assume a *processor* has 8 cores, and a multi-processor *server* 4 processors; then a server did 2^{38} cycles/second.
-
- In a year there are approx. 2^{25} seconds. A *server-year* (s-y) does 2^{63} cycles, so collisions after (74-63): 2^{11} s-y.

Hashing



- How insecure is SHA-1?
 - Schneier reports that in approx. 2^{74} computer cycles a SHA-1 collision is found
 - In 2016 a *core* ran at approx. 2^{33} cycles/sec. Assume a *processor* typically has 8 cores, and a multi-processor *server* 4 processors; then a server did 2^{38} cycles/second.
 - In a year there are approx. 2^{25} seconds. A *server-year (s-y)* does 2^{63} cycles, so *collisions after (74-63): 2^{11} s-y.*
-
- Using *Moore's law* (computing power doubles every 18 months):
 - In **2019**: $3/1.5 = 2$ doublings in computer power (2^2): 2^{65} cycles, so *collisions after: 2^9 s-y.*
 - In **2022**: $6/1.5 = 4$ doublings (2^4): 2^{67} cycles, *collisions after: 2^7 s-y.*


Hashing



- How insecure is SHA-1?
 - *Schneier* reports that in approx. 2^{74} computer cycles a SHA-1 collision is found
 - A core runs at approx. 2^{33} cycles/sec. In 2016 processors typically had 8 cores, a multi-processor server had 4 processors, so a server did 2^{38} cycles/second.
 - In a year there are approx. 2^{25} seconds. A server-year (s-y) did 2^{63} cycles, collision after: 2^{11} s-y.
 - Using *Moore's law* (computing power doubles every 18 months):
 - In 2019: $3/1.5 = 2$ doublings in computer power (2^2): 2^{65} cycles, so collisions after: 2^9 s-y.
 - In 2022: $6/1.5 = 4$ doublings (2^4): 2^{67} cycles, collisions after: 2^7 s-y.
-
- Renting a server costs approx. €250/yr = approx. € 2^8 /yr, multiply by #s-y for a collision: a collision attack in 2016 costs approx. € 2^{19} , approx. €500k, in 2019 2^{17} (€130k), in 2022 2^{15} (€33k).

Sharing Secrets

- Simple ways to share a secret
 - Basic idea: *polynomials fitting*
 - E.g., straight line - polynomial of degree 1
 - Given two points, the line's equation can be determined



How do we calculate Y_0 ?

$$Y_0 = Y_1 - X_1 * (Y_2 - Y_1) / (X_2 - X_1)$$

- In general: $n + 1$ points are required to determine a polynomial of degree n .

Sharing Secrets


- Simple ways to share a secret
 - Select a polynomial of your choice
 - E.g., $Y = aX + b$
 - A polynomial of degree 1: Alice and Bob each receive one point on this line as the secret info.
 - Using only their own point neither Alice nor Bob can determine the secret.
 - The secret could be, e.g., the Y coordinate for $X = 0$

Sharing Secrets

- Simple ways to share a secret
 - E.g., $Y = aX + b$
 - Two parties each receive the coordinates of one point on this line.
 - The secret S is defined as point $(0, S)$, so $b == S$
 - Select a and two points X_a and X_b .
 - Alice gets (X_a, Y_a) , where $Y_a = a * X_a + S$,
Bob gets (X_b, Y_b) , where $Y_b = a * X_b + S$.

Key Escrow

- An example use of secret sharing
- Key Escrow
 - Store your secret (i.e., key) with a *trusted* party.
 - A notary? A good friend?



Jobs in which nobody understands what you do.



- Who do **you** trust??

Key Escrow

- Key Escrow
 - Example: The *clipper chip*, announced in 1993 and abandoned by 1996.
 - To be built into all electronic devices offering cryptography
 - Used a symmetric encryption algorithm (*Skipjack*) comparable to DSA
 - *Key escrow* by the US Government...
 - Can the Government be trusted?



Key Escrow

- Key Escrow
 - Alternative:
 - use polynomial key-splitting, requiring n people to work together to determine your secret, which may be the *passphrase* to unlock your *file of secrets* or to access your *encrypted file system*.

Key Escrow

- Key Escrow
 - Subtle modification:
 - A polynomial of order $n-1$ may be determined if n points are provided.
 - Provide m points ($m > n$) to m people, thus implementing an

n out of m

key escrow: any n people may join to obtain the secret.

Steganography

- Information Hiding (steganography)
 - Hide information in unlikely places
 - *Yes, the alchemists worshipped the antimatter, vowing dark, agile actions while announcing algebra...*
 - The problem is of course Kerckhoffs principle
 - Unused places can be used to hide information in
 - cf. Stamp's low order bits of a html-file's color attribute
 - *Collusion attacks* (i.e., use *diff*) can be used to reveal hidden information.

What did we learn today?

- Topics this lecture:
 - Concept and requirements
 - Collisions
 - Use of hashing in cryptography
 - CRC: a non-cryptographic hash
 - MD5, SHA-x, Tiger
 - Sharing Secrets / Key Escrow
 - Information Hiding
 - E-mail peculiarities

FAQ

- What is m in Tiger hash?
 - It is a constant value (denoting the round index, i.e., 5,7,9)
- Does Tiger hash work with chaining, meaning for instance the resulting a, b, c , values of hashing X_0 would be used in hashing X_1 , given $M = \{X_0, \dots, X_n\}$?
 - Yes, the result of the hashing of X_i is the final hash value or the initial value for the next message block X_{i+1} .
- Why does $h(M, K)$ prevent length extension attacks?
 - TBC

That's all for today.