



university of  
 groningen

# Computer Architecture 2023-24 (WBCSo10-05)

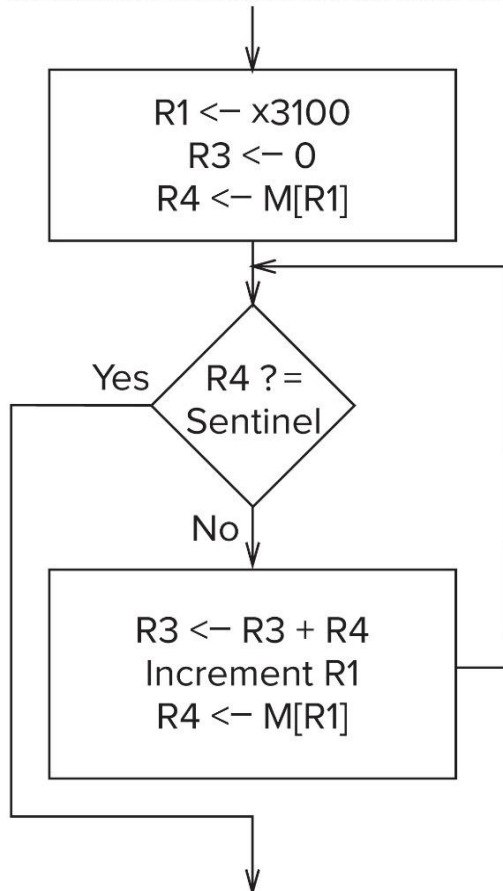
## Lecture 7: The LC-3 (Chapters 5)

Reza Hassanpour  
r.zare.hassanpour@rug.nl

# Example: Loop with Sentinel

- Compute the sum of integers starting at x3100, until a negative integer is found. Instructions start at x3000

Copyright © McGraw-Hill Education. Permission required for reproduction or display.




Address	15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0	
x3000	1	1	1	0	0	0	1	0	1	1	1	1	1	1	1	1	R1 ← x3100
x3001	0	1	0	1	0	1	1	0	1	1	1	0	0	0	0	0	R3 ← 0
x3002	0	1	1	0	1	0	0	0	0	1	0	0	0	0	0	0	R4 ← M[R1]
x3003	0	0	0	0	1	0	0	0	0	0	0	0	0	1	0	0	BRn x3008
x3004	0	0	0	1	0	1	1	0	1	1	0	0	0	1	0	0	R3 ← R3+R4
x3005	0	0	0	1	0	0	1	0	0	1	1	0	0	0	0	1	R1 ← R1+1
x3006	0	1	1	0	1	0	0	0	0	1	0	0	0	0	0	0	R4 ← M[R1]
x3007	0	0	0	0	1	1	1	1	1	1	1	1	0	1	1		BRnzp x3003

- In this case, it's the data that tells when to exit the loop. When the value loaded into R4 is negative, take branch to x3008

# Unconditional Branch (JMP)

- › We can create a branch that is always taken by setting n, z, and p to 1
- › However, the target address of BR is limited by the 9-bit PC offset
- › The JMP instruction provides an unconditional branch to any target location by using the contents of a register as the target address. It simply copies the register into the PC



# TRAP: Invoke a System Service Routine

- › The TRAP instruction is used to give control to the operating system to perform a task that user code is not allowed to do. The details will be explained in Chapter 9



- › For now, you just need to know that bits [7:0] hold a "trap vector" -- a unique code that specifies the service routine. The service routines used in this part of the course are:

trapvector	service routine
x23	Input a character from the keyboard
x21	Output a character to the monitor
x25	Halt the processor

# Input / Output Service Routines

- › Getting character input from the keyboard
  - TRAP x23 is used to invoke the keyboard input service routine.
  - When the OS returns control to our program, the **ASCII code** for the key pressed by the user will be **in Ro**.
  
- › Sending character output to the monitor
  - TRAP x21 is used to invoke the monitor output service routine.
  - **Before** invoking the routine, put the **ASCII character** to be output **into Ro**.

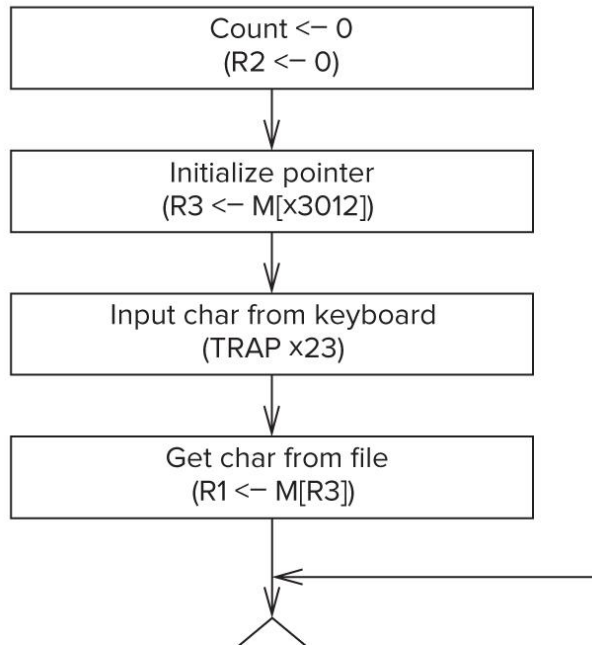
# Example Program: Count Occurrences of a Character

- › We want to count the number of times a user-specified character appears in a text file, and then print the count to the monitor
- The text file is stored in memory as a sequence of ASCII characters
- The end of the file is denoted with the **EOT**\* character (x04).  
NOTE: We do not know how many characters are in the file; EOT is the sentinel value that signals when we are done
- **A pointer to the file** will be stored at the end of the program. (A "pointer" is a memory address; it will be the address of the first character in the file)
- We will assume that the character will appear **no more than 9** times
- Program instructions will start at **x3000**. The **file data** can be anywhere in memory

\* *End Of Transmission*

# Part 1: Initializing Registers

Copyright © McGraw-Hill Education. Permission required for reproduction or display.



Copyright © McGraw-Hill Education. Permission required for reproduction or display.

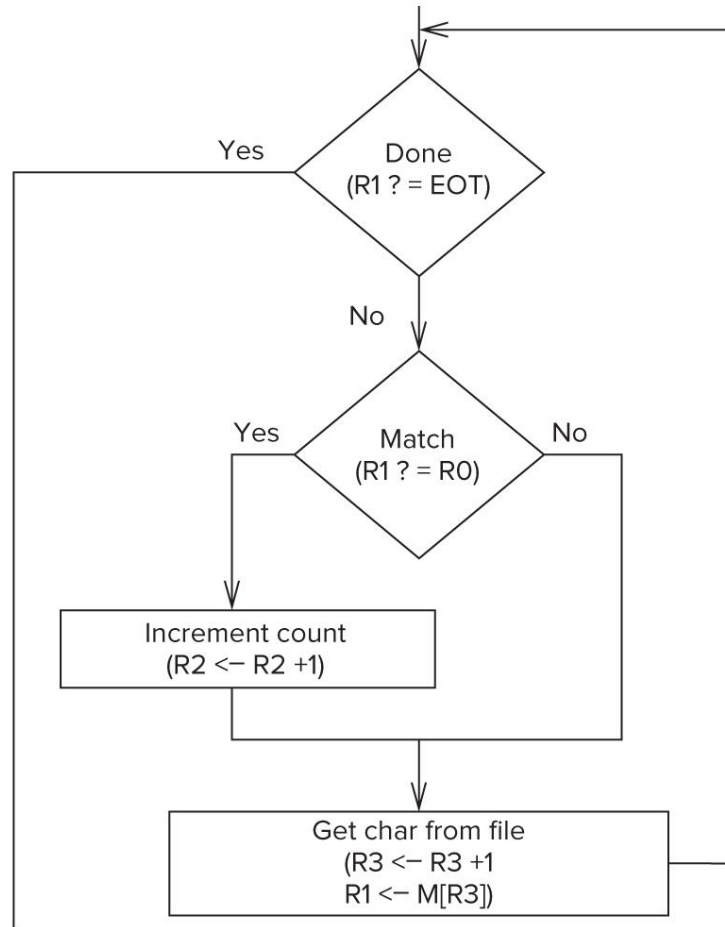
Address	15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0	
x3000	0	1	0	1	0	1	0	0	1	0	1	0	0	0	0	0	R2 ← 0
x3001	0	0	1	0	0	1	1	0	0	0	0	1	0	0	0	0	R3 ← M[x3012]
x3002	1	1	1	1	0	0	0	0	0	0	1	0	0	0	1	1	TRAP x23
x3003	0	1	1	0	0	0	1	0	1	1	0	0	0	0	0	0	R1 ← M[R3]
x3004	0	0	0	1	1	0	0	0	0	1	1	1	1	0	0	0	R4 ← R1-4
x3005	0	0	0	0	0	1	0	0	0	0	0	0	1	0	0	0	BRz x300E
x3006	1	0	0	1	0	0	1	0	0	1	1	1	1	1	1	1	R1 ← NOT R1
x3007	0	0	0	1	0	0	1	0	0	1	1	0	0	0	0	1	R1 ← R1 + 1
x3008	0	0	0	1	0	0	1	0	0	1	0	0	0	0	0	0	R1 ← R1 + R0
x3009	0	0	0	0	1	0	1	0	0	0	0	0	0	0	0	1	BRnp x300B
x300A	0	0	0	1	0	1	0	0	1	0	1	0	0	0	0	1	R2 ← R2 + 1
x300B	0	0	0	1	0	1	1	0	1	1	1	0	0	0	0	1	R3 ← R3 + 1
x300C	0	1	1	0	0	0	1	0	1	1	0	0	0	0	0	0	R1 ← M[R3]
x300D	0	0	0	0	1	1	1	1	1	1	1	1	0	1	1	0	BRnzp x3004
x300E	0	0	1	0	0	0	0	0	0	0	0	0	0	1	0	0	R0 ← M[x3013]
x300F	0	0	0	1	0	0	0	0	0	0	0	0	0	0	1	0	R0 ← R0 + R2
x3010	1	1	1	1	0	0	0	0	0	0	1	0	0	0	0	1	TRAP x21
x3011	1	1	1	1	0	0	0	0	0	0	1	0	0	1	0	1	TRAP x25
x3012	Starting address of file																
x3013	0	0	0	0	0	0	0	0	0	0	1	1	0	0	0	0	ASCII TEMPLATE

R2 is counter. R3 is address of first character to read from file.  
R0 is character from keyboard. R1 is first character from file.



# Part 2: Read Characters and Count

Copyright © McGraw-Hill Education. Permission required for reproduction or display.



Copyright © McGraw-Hill Education. Permission required for reproduction or display.


Address	15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0	
x3000	0	1	0	1	0	1	0	0	1	0	1	0	0	0	0	0	R2 <- 0
x3001	0	0	1	0	0	1	1	0	0	0	0	1	0	0	0	0	R3 <- M[x3012]
x3002	1	1	1	1	0	0	0	0	0	0	1	0	0	0	1	1	TRAP x23
x3003	0	1	1	0	0	0	1	0	1	1	0	0	0	0	0	0	R1 <- M[R3]
x3004	0	0	0	1	1	0	0	0	0	1	1	1	1	1	0	0	R4 <- R1-4
x3005	0	0	0	0	0	1	0	0	0	0	0	0	1	0	0	0	BRz x300E
x3006	1	0	0	1	0	0	1	0	0	1	1	1	1	1	1	1	R1 <- NOT R1
x3007	0	0	0	1	0	0	1	0	0	1	1	0	0	0	0	1	R1 <- R1 + 1
x3008	0	0	0	1	0	0	1	0	0	1	0	0	0	0	0	0	R1 <- R1 + R0
x3009	0	0	0	0	1	0	1	0	0	0	0	0	0	0	0	1	BRnp x300B
x300A	0	0	0	1	0	1	0	0	1	0	1	0	0	0	0	1	R2 <- R2 + 1
x300B	0	0	0	1	0	1	1	0	1	1	1	0	0	0	0	1	R3 <- R3 + 1
x300C	0	1	1	0	0	0	1	0	1	1	0	0	0	0	0	0	R1 <- M[R3]
x300D	0	0	0	0	1	1	1	1	1	1	1	1	0	1	1	0	BRnzp x3004
x300E	0	0	1	0	0	0	0	0	0	0	0	0	0	1	0	0	R0 <- M[x3013]
x300F	0	0	0	1	0	0	0	0	0	0	0	0	0	0	1	0	R0 <- R0 + R2
x3010	1	1	1	1	0	0	0	0	0	0	1	0	0	0	0	1	TRAP x21
x3011	1	1	1	1	0	0	0	0	0	0	1	0	0	1	0	1	TRAP x25
x3012	Starting address of file																
x3013	0	0	0	0	0	0	0	0	0	0	1	1	0	0	0	0	ASCII TEMPLATE

Compare character to immediate x04 (EOT in ASCII). If equal, exit loop. Otherwise, count if matches user input and read the next character.



# Part 3: Output Count and HALT

Copyright © McGraw-Hill Education. Permission required for reproduction or display.



Copyright © McGraw-Hill Education. Permission required for reproduction or display.

Address	15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0	
x3000	0	1	0	1	0	1	0	0	1	0	1	0	0	0	0	0	R2 <- 0
x3001	0	0	1	0	0	1	1	0	0	0	0	1	0	0	0	0	R3 <- M[x3012]
x3002	1	1	1	1	0	0	0	0	0	0	1	0	0	0	1	1	TRAP x23
x3003	0	1	1	0	0	0	1	0	1	1	0	0	0	0	0	0	R1 <- M[R3]
x3004	0	0	0	1	1	0	0	0	0	1	1	1	1	1	0	0	R4 <- R1-4
x3005	0	0	0	0	0	1	0	0	0	0	0	0	1	0	0	0	BRz x300E
x3006	1	0	0	1	0	0	1	0	0	1	1	1	1	1	1	1	R1 <- NOT R1
x3007	0	0	0	1	0	0	1	0	0	1	1	0	0	0	0	1	R1 <- R1 + 1
x3008	0	0	0	1	0	0	1	0	0	1	0	0	0	0	0	0	R1 <- R1 + R0
x3009	0	0	0	0	1	0	1	0	0	0	0	0	0	0	0	1	BRnp x300B
x300A	0	0	0	1	0	1	0	0	1	0	1	0	0	0	0	1	R2 <- R2 + 1
x300B	0	0	0	1	0	1	1	0	1	1	1	0	0	0	0	1	R3 <- R3 + 1
x300C	0	1	1	0	0	0	1	0	1	1	0	0	0	0	0	0	R1 <- M[R3]
x300D	0	0	0	0	1	1	1	1	1	1	1	1	0	1	1	0	BRnzp x3004
x300E	0	0	1	0	0	0	0	0	0	0	0	0	0	1	0	0	R0 <- M[x3013]
x300F	0	0	0	1	0	0	0	0	0	0	0	0	0	0	1	0	R0 <- R0 + R2
x3010	1	1	1	1	0	0	0	0	0	0	1	0	0	0	0	1	TRAP x21
x3011	1	1	1	1	0	0	0	0	0	0	1	0	0	1	0	1	TRAP x25
x3012	Starting address of file																
x3013	0	0	0	0	0	0	0	0	0	0	1	1	0	0	0	0	ASCII TEMPLATE

When loop is finished, convert count (R2) to the corresponding ASCII character by adding 'o' (x30). Output the character and halt the program

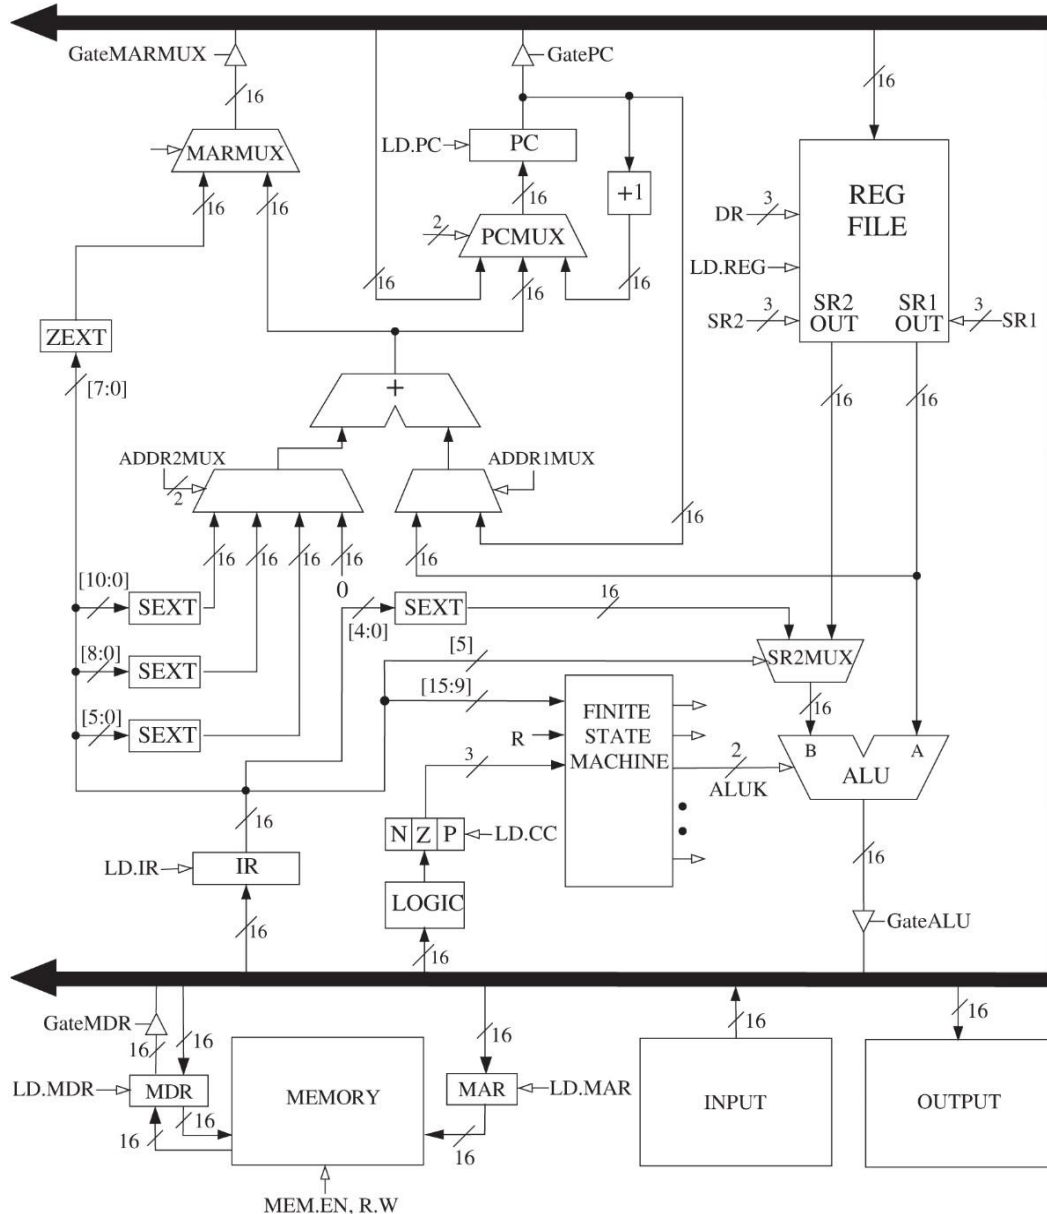


# ASCII Codes

Hex	Value	Hex	Value	Hex	Value	Hex	Value	Hex	Value	Hex	Value	Hex	Value	Hex	Value
00	NUL	10	DLE	20	SP	30	0	40	@	50	P	60	`	70	p
01	SOH	11	DC1	21	!	31	1	41	A	51	Q	61	a	71	q
02	STX	12	DC2	22	"	32	2	42	B	52	R	62	b	72	r
03	ETX	13	DC3	23	#	33	3	43	C	53	S	63	c	73	s
04	EOT	14	DC4	24	\$	34	4	44	D	54	T	64	d	74	t
05	ENQ	15	NAK	25	%	35	5	45	E	55	U	65	e	75	u
06	ACK	16	SYN	26	&	36	6	46	F	56	V	66	f	76	v
07	BEL	17	ETB	27	'	37	7	47	G	57	W	67	g	77	w
08	BS	18	CAN	28	(	38	8	48	H	58	X	68	h	78	x
09	HT	19	EM	29	)	39	9	49	I	59	Y	69	i	79	y
0A	LF	1A	SUB	2A	*	3A	:	4A	J	5A	Z	6A	j	7A	z
0B	VT	1B	ESC	2B	+	3B	;	4B	K	5B	[	6B	k	7B	{
0C	FF	1C	FS	2C	,	3C	<	4C	L	5C	\	6C	l	7C	
0D	CR	1D	GS	2D	-	3D	=	4D	M	5D	]	6D	m	7D	}
0E	SO	1E	RS	2E	.	3E	>	4E	N	5E	^	6E	n	7E	~
0F	SI	1F	US	2F	/	3F	?	4F	O	5F	_	6F	o	7F	DEL

# LC-3 Data Path 1

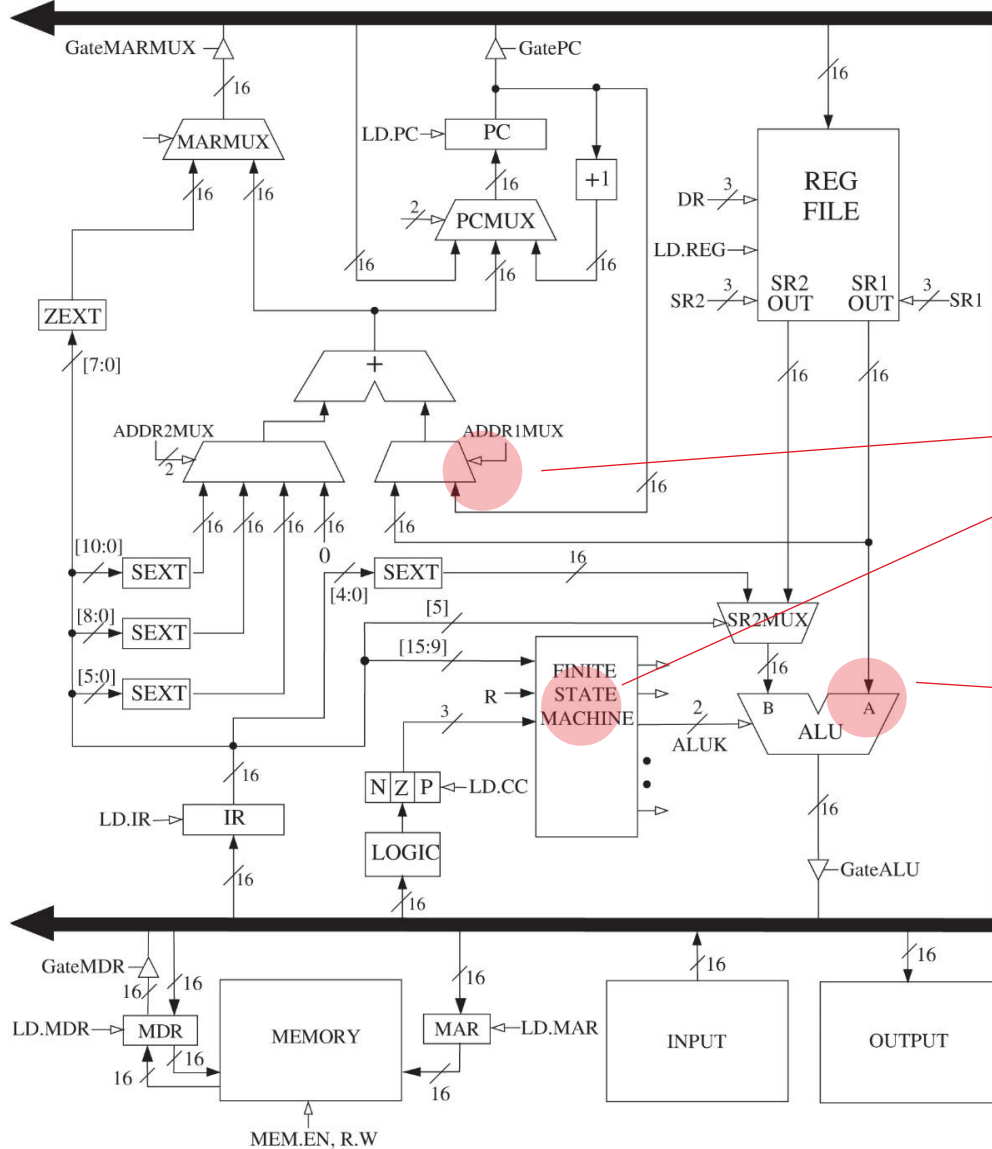
Copyright © McGraw-Hill Education. Permission required for reproduction or display.



- › Data path is used to execute LC-3 programs
- › PC is initialized to point to the first instruction. Clock is enabled, and the control unit takes over
- › Next slides will give a little more detail on various components

## LC-3 Data Path 2

Copyright © McGraw-Hill Education. Permission required for reproduction or display.

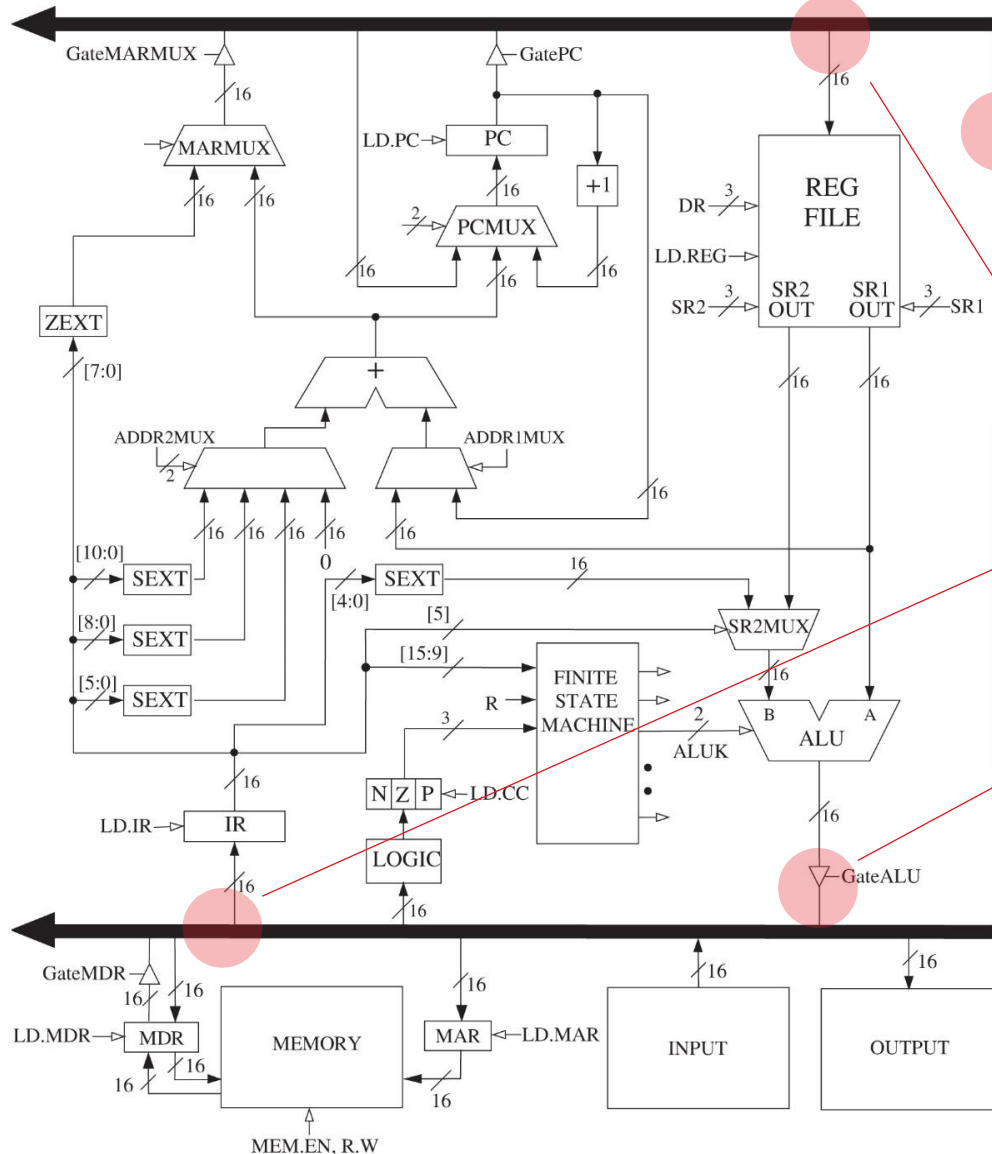


Arrows with open heads  
represent control signals  
from FSM

Arrows with filled heads  
represent data that is processed

# LC-3 Data Path 3

Copyright © McGraw-Hill Education. Permission required for reproduction or display.



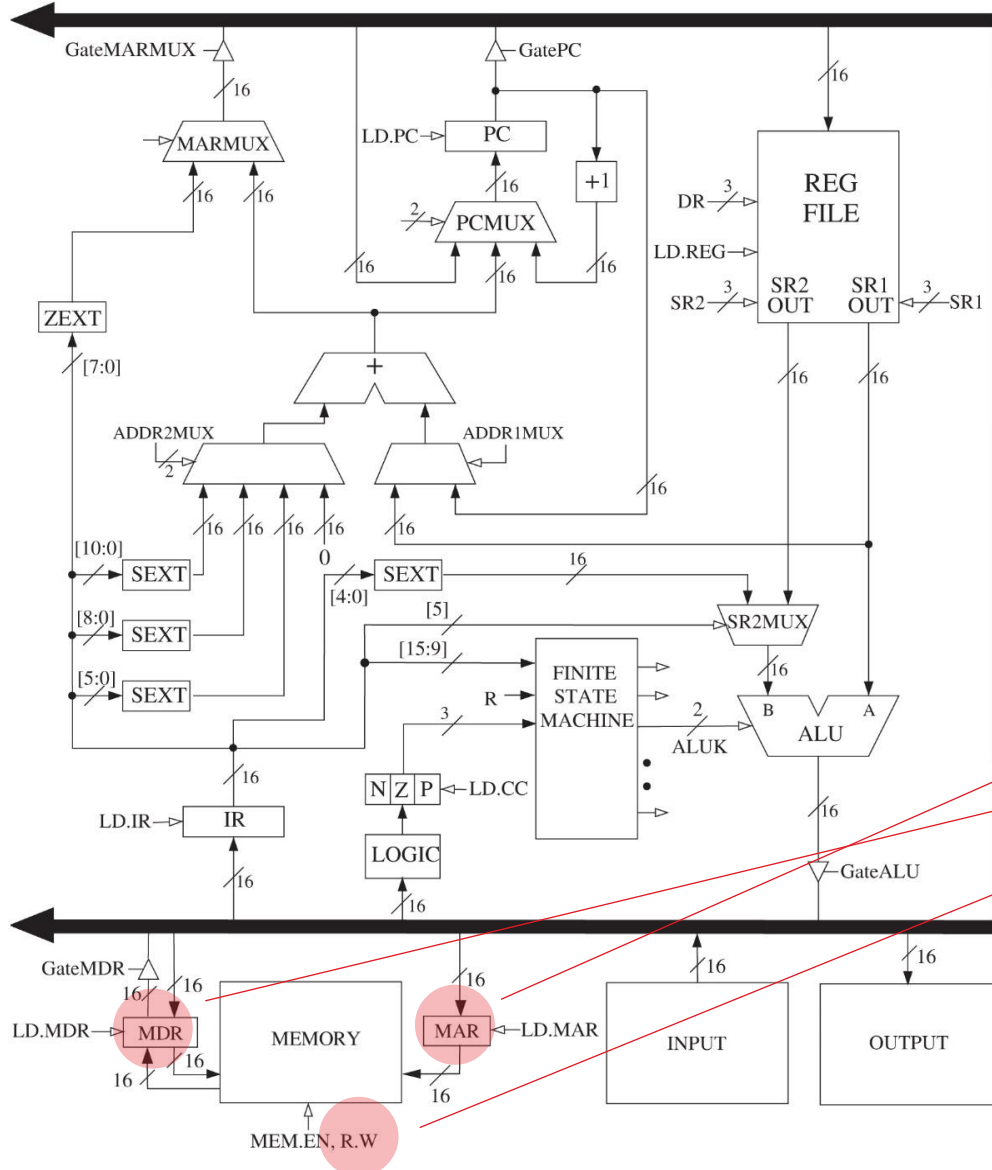
**Global bus** is a set of wires that allow various components to transfer 16-bit data to other components.

One or more components may read data from the bus on any cycle.

Tri-state device determines which component puts data on the bus. Only one source of data at any time.

# LC-3 Data Path 4

Copyright © McGraw-Hill Education. Permission required for reproduction or display.

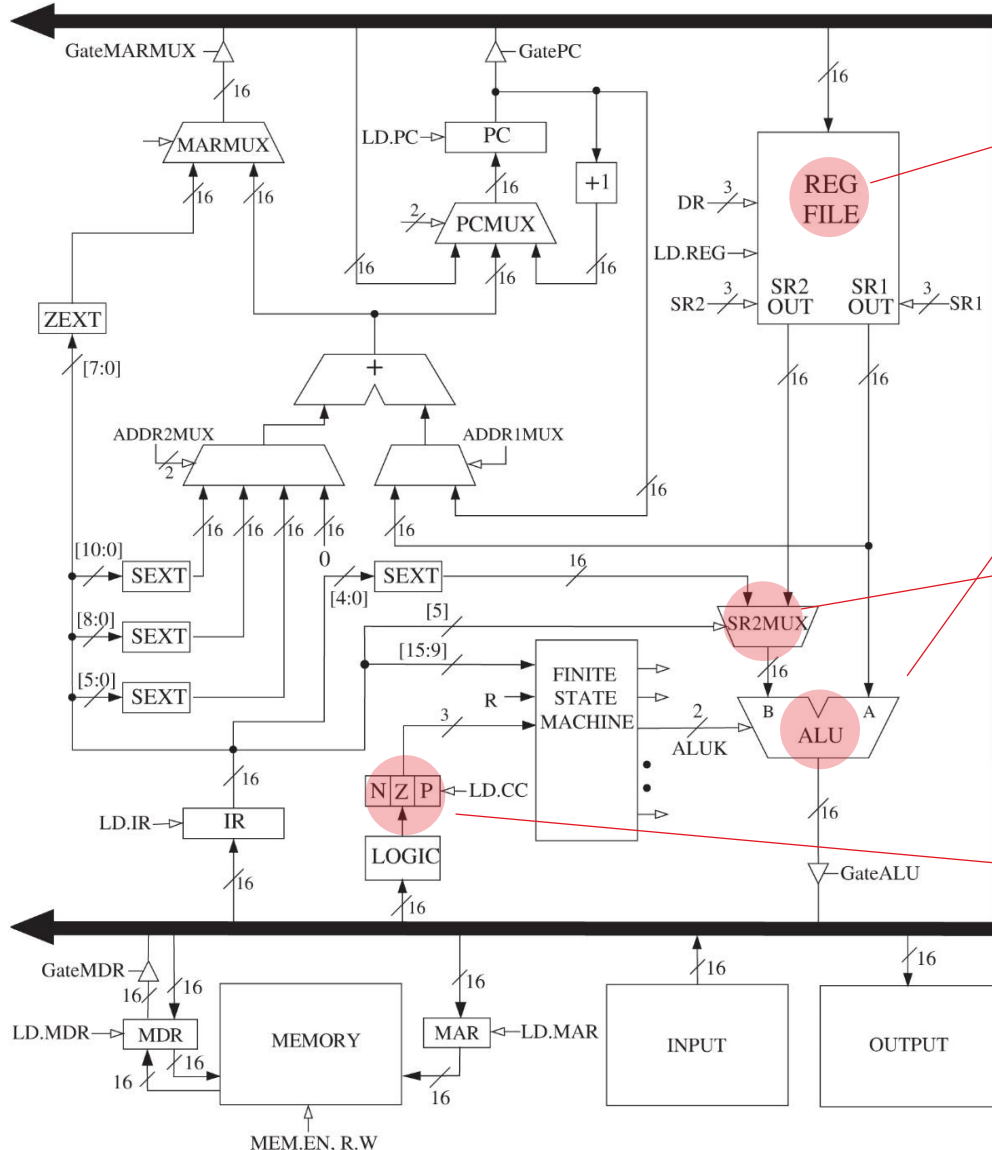


Memory interface:  
MAR  
MDR  
Read/Write  
control



# LC-3 Data Path 5

Copyright © McGraw-Hill Education. Permission required for reproduction or display.



## Register File (R0-R7)

Control signals specify two source register (SR1, SR2) and one destination (DR).

ALU performs ADD, AND, NOT.

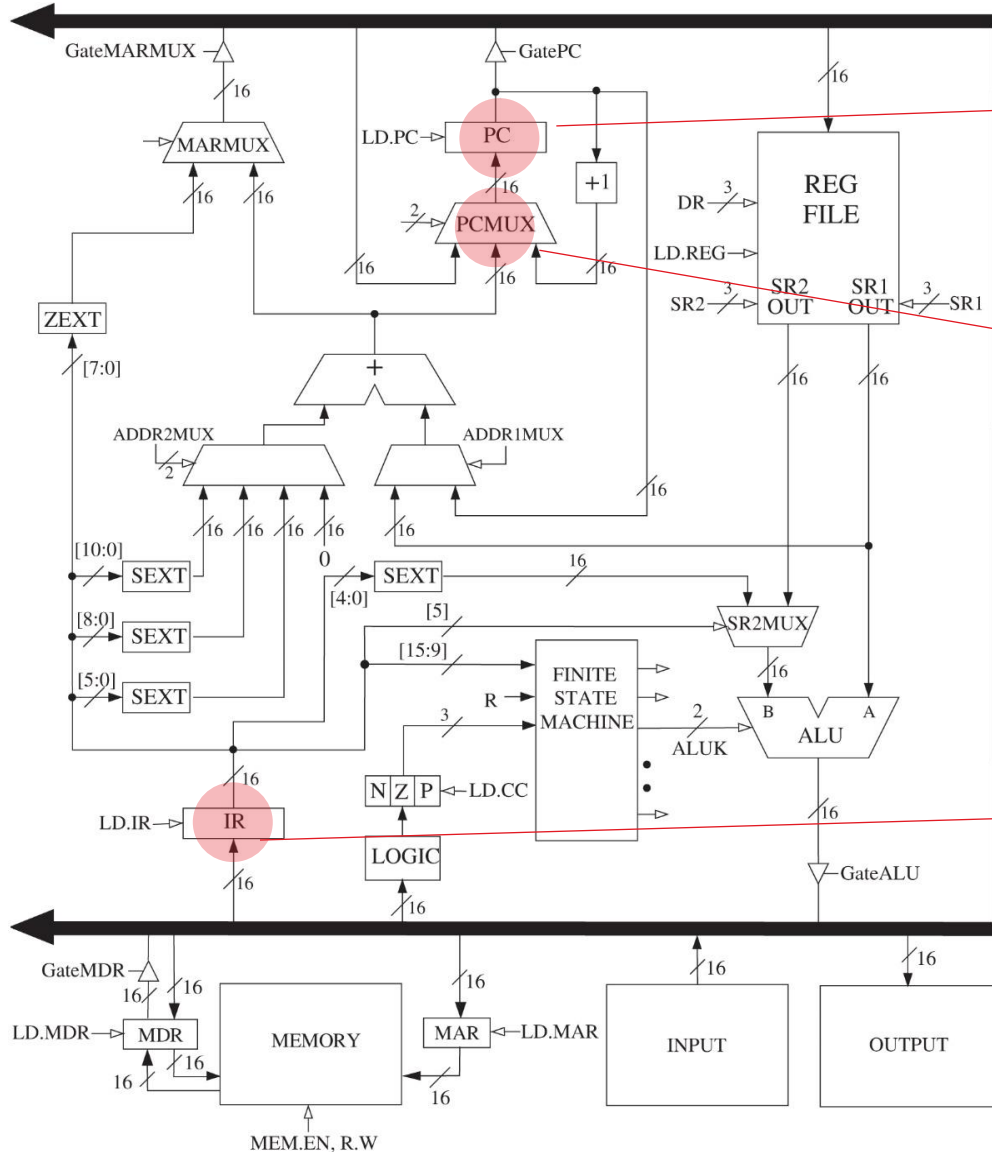
Operand A always comes from register file. Operand B is from register file or IR. Output goes to bus, to be written into register file.

Condition codes are set by looking at data placed on the bus by ALU or memory (MDR).



# LC-3 Data Path 6

Copyright © McGraw-Hill Education. Permission required for reproduction or display.



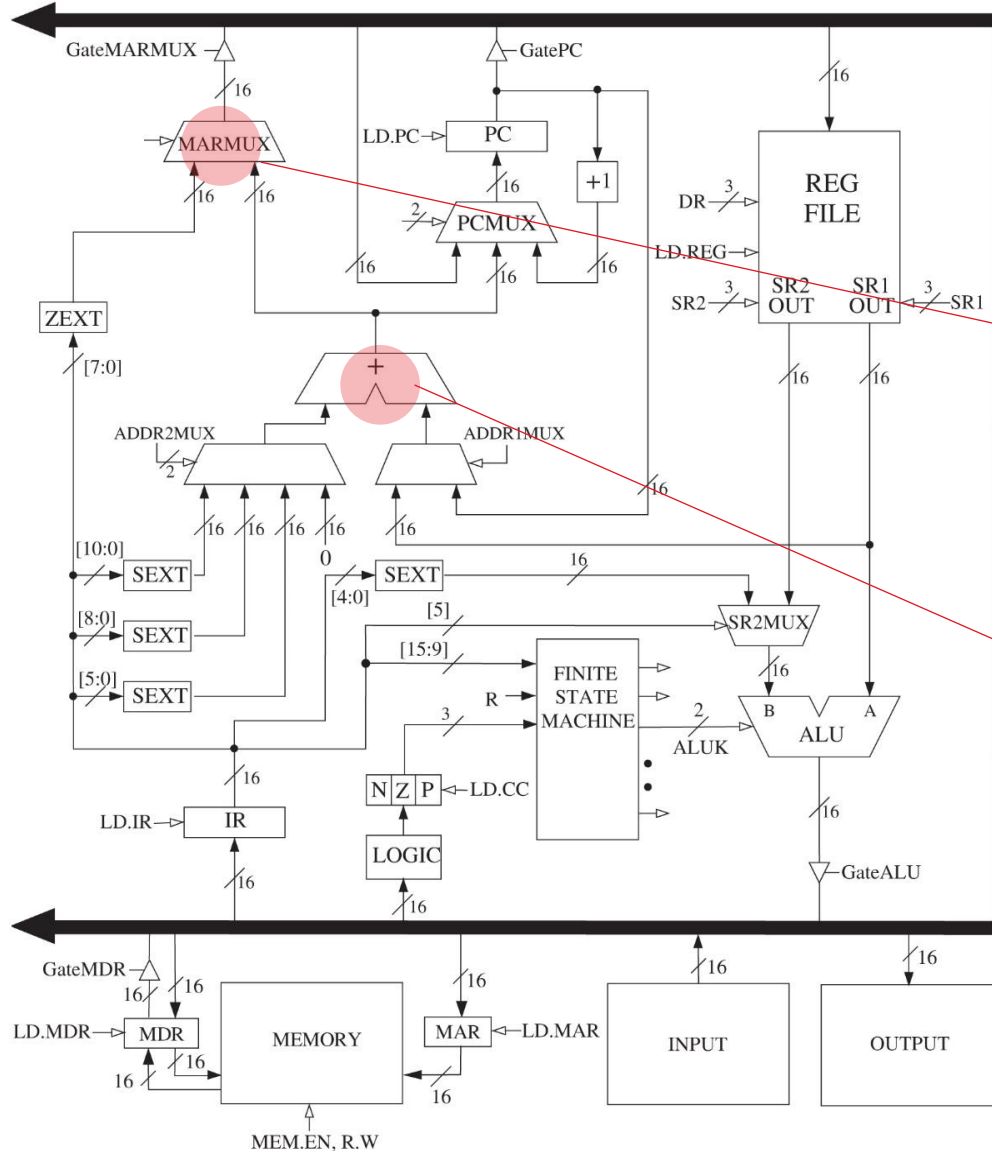
**PC** puts address on bus.  
Placed in MAR during Fetch

**PCMUX** allows various  
values to be written to PC:  
incremented PC (Fetch),  
computed address (BR), or  
register data from bus (JMP)

**IR** gets data from bus (MDR)  
during Fetch

# LC-3 Data Path 7

Copyright © McGraw-Hill Education. Permission required for reproduction or display.



**MARMUX** chooses value to be written to MAR during load, store, or TRAP

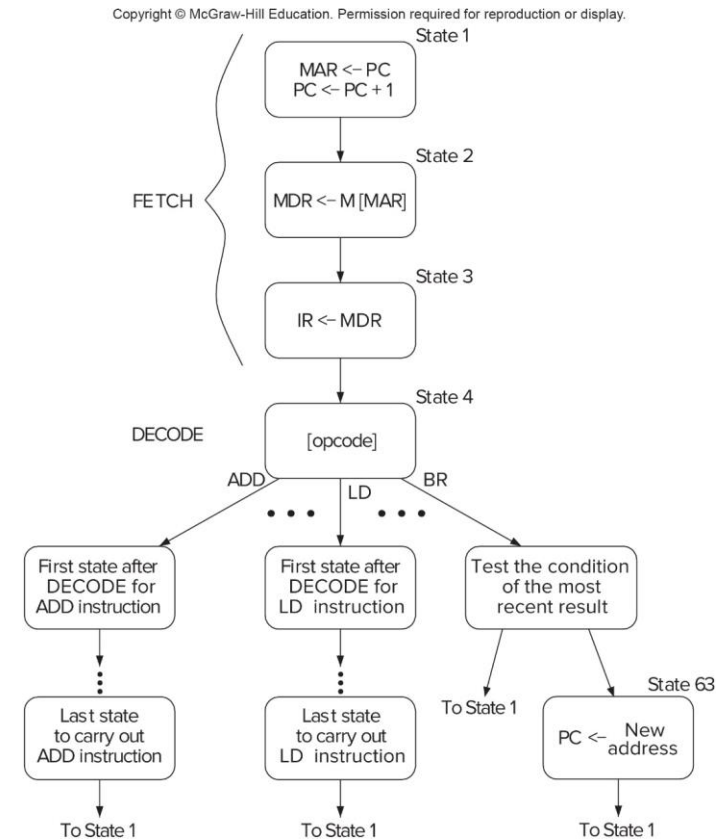
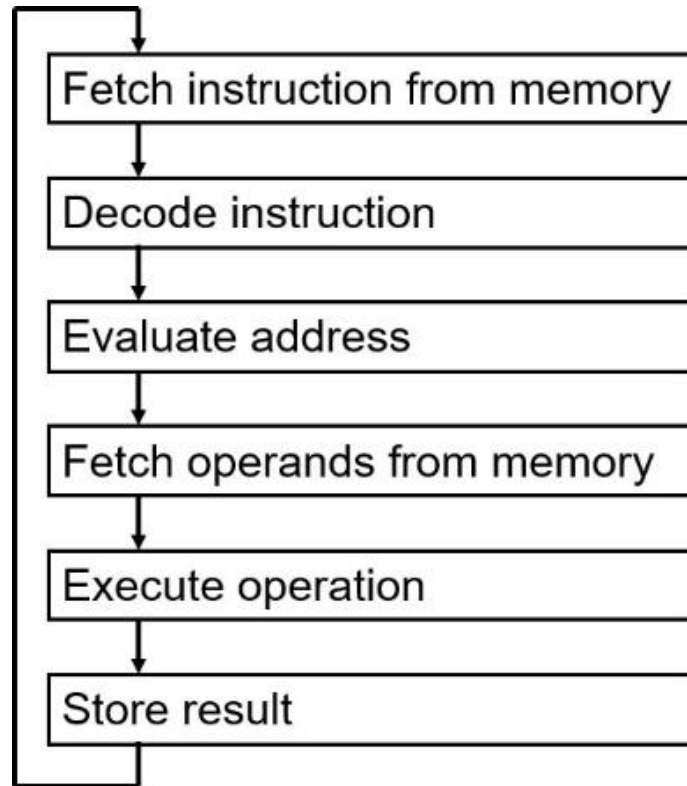
**Evaluate Address** phase adds offset to PC or register for load, store, BR

# Optimizing the Performance

- › Memory access (Load/Store) operations are slow compared to the execution time in the CPU.
- › For better performance,
  - Access memory in advance
  - Use faster memories


The remaining slides of this lecture are not from the test books

# Machine cycle is not instruction cycle




$T$  (period) [s]  
 $f$  (frequency) [Hz]  
 $T = 1/f$   
 $2 \text{ ns} \rightarrow 500 \text{ MHz}$

Copyright © McGraw-Hill Education. Permission required for reproduction or display.



Copyright © McGraw-Hill Education. Permission required for reproduction or display.



# On instruction complexity

- Originally computers had extremely basic instructions
- Later more complex e.g., *floating point* instructions
- Ever bigger sets of hardware-based instructions due to:
  - Instruction compatibility requirements
  - Rising cost of software development
- **Interpreted instruction sets** on low cost computers
  - Easier to fix, ability to add new instructions, efficient development of complex instructions

# RISC vs CISC

- **pre 1980s:** Complex instructions were used in an attempt to bridge the “semantic gap”
- **in 1980s:** RISC (Reduced vs Complex Instruction Set Computer) counter movement
- Emphasis on **faster issuing** of instructions
- **1990s and onward:** hybrid models starting from Intel’s 486 architecture
  - Common instructions are fast in RISC sense
  - Complex (but uncommon) instructions are slow but easier to implement

## (RISC) design principles for modern computers

- All instructions must be directly **executed by hardware**
- **Maximize the rate** at which instructions are issued
- Instructions should be **easy to decode**
- Only **load and store** instructions should reference memory
- Provide **plenty of registers**



# On parallelism

- **Instruction-level**

- More instructions per second
- **Pipelining** and **Superscalar architectures**


- **Processor-level**

- Hard limit of information transfer ( $\sim 20$  cm/nsec)
- Faster, larger chips ---> more heat
- The number of transistors per unit area still increases
- Adding more CPUs on a chip is the only way to improve
- (*back to*) **Data Parallel Computers**, **Multiprocessors**, and **Multicomputers** (what about Cloud, e.g., AWS and Azure?)

# Pipelining

- **Prefetch buffers** already available since 1959  
→ fetching & execution stages
- A **pipeline** divides instruction execution into many **stages** running in parallel, synchronised by the CPU clock
- Trade-off between latency and processor bandwidth (see next slide)

What can be observed about these timings:


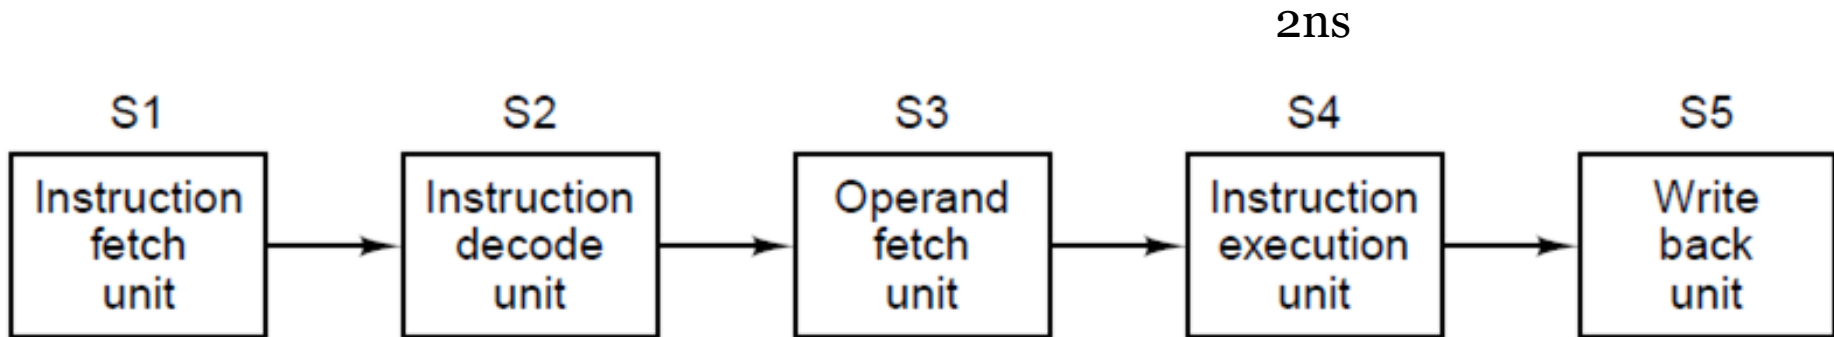


Assume the following implementation of a 5-stage instruction cycle

# Pipelining metrics

- Cycle time  $T$  (in ns), number of stages  $n$
- **Latency** =  $n * T$  (in ns)
  - Execution time for each instruction in nanoseconds
- **Bandwidth** =  $1/T$  (expressed in Million IPS: Instructions per second)
  - How many MIPS the CPU is capable of delivering

# Pipelining example (5 stages)




- $T = 2 \text{ ns}, n = 5$
- Latency =  $n * T = 10 \text{ ns}$
- Bandwidth =  $1/T = 1 / 2 * 10^{-9} = 500 * 10^6 \text{ IPS} = 500 \text{ MIPS}$

2 ns → 500MHz

Intel Pentium Pro:  
541 MIPS at 200 MHz  
(1996)

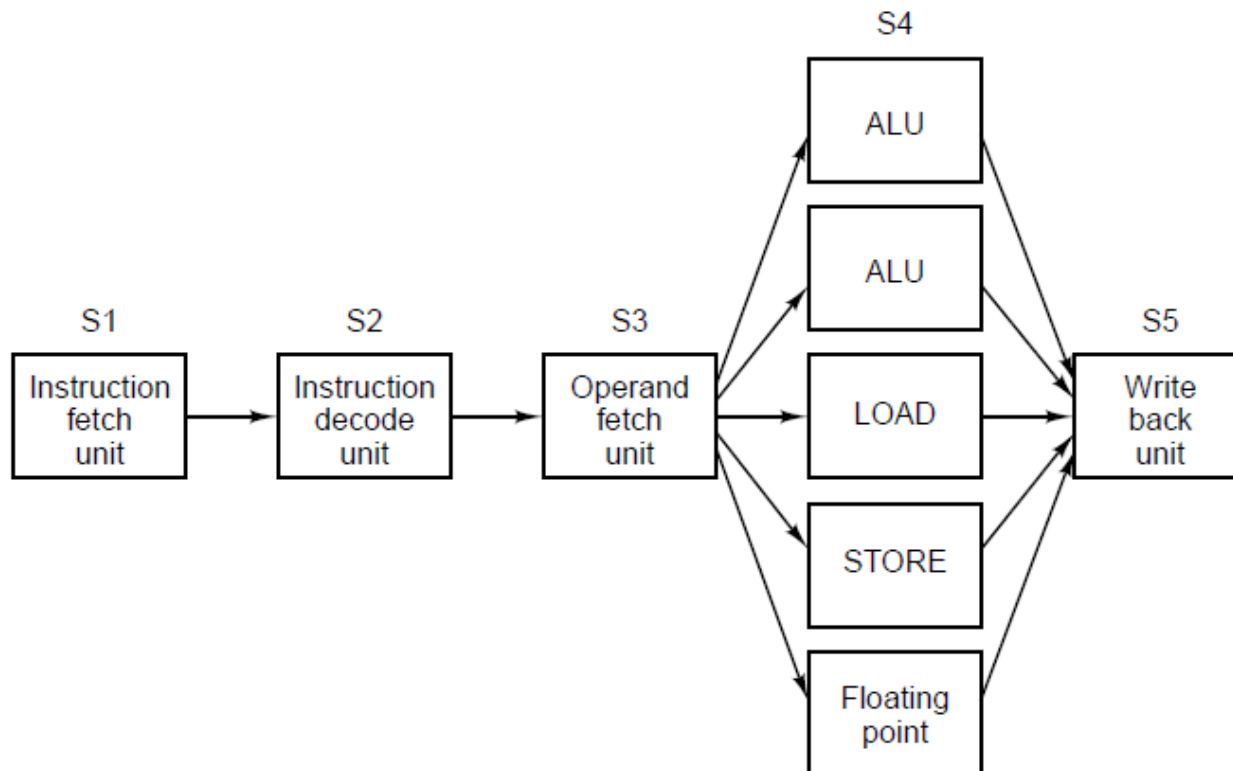
# Multiple pipelines



- Conflict prevention happens at the compiler (software) level
- Different purposes of pipelines also possible
  - Intel's Pentium architecture: u/v (regular/integer arithmetic-specific) pipelines

# Superscalar architectures

- **Key observation:** S4 (instruction execution) may take much longer than S3 (operand fetching) → multiple execution units can speed up execution time



# What about cache memories

- **Main issue with computer Main memories:** designed for capacity, **not for speed**
  - CPUs are getting faster at a faster pace
  - More cores on the CPU means more Main Memory refs
  - Moving (more) memory inside the CPU increases construction cost




# Cache memory motivation

- Possible solutions

- 1) Issue **read to Main memory and continue execution** until word is needed; then **stall** (too many wasted cycles)
- 2) Issue **reads in advance** through compiler optimization (ends up in a **software stall**)
- 3) Add intermediate memory level of fast (=expensive) memory as close to CPU as possible – called **cache**

# Cache memory basics

- Most heavily used memory words to be kept in cache
- **Principle of locality**
  - If address  $A$  is read, then *most probably*  $A+1$  will be read in the next time interval, or
  - Memory reads in a short time interval only use a small part of the (total) memory
- **Fixed size blocks** are transferred between cache and memory (called cache lines)



# Cache quality

- Access times:  $c$  for the cache,  $m$  for main memory
  - Refer  $k$  times to address A: 1(st) reference to main memory,  $k-1$  references to the cache
  - The cache hit ratio  $h$  is the proportion of cache references:
$$h = (k-1)/k$$
  - Average access time:  $(1-h)m + c$
  - As  $k$  increases,  $h \rightarrow 1$  and  $m$  “vanishes” from the equation



# Cache Access Methods

- › How can we decide if a data block is in cache or not?
- › Because there are fewer cache lines than main memory blocks, an algorithm is needed for mapping main memory blocks into cache lines.
- › There are three methods to access cache data:
  - **Direct** Mapping
  - **Associative** Mapping
  - **Set Associative** Mapping



# Direct Mapping

- › It is the simplex technique, maps each block of main memory into only one possible cache line i.e. a given main memory block can be placed in one and only one place on cache.
- ›  $i = j \text{ modulo } m$ ,
- › Where  $i$  = cache line number;
- ›  $j$  = main memory block number;
- ›  $m$  = number of lines in the cache



# Example

- › Assume memory is 64K Words of 4 bytes,
- › 32 bit addressability, 16 address lines
- › Cache is 512 words (512x32)
- › To show each cache row we need 9 bits
- › Remaining 7 bits of the address of the word in memory is used as tag



# Direct Mapping Cache

	Data
x0000	1234
x0001	2735
x0002	7428
x0003	9082
x0402	0012

Memory

Tag (7bits)	Data
x00	1234
x02	0012

Cache





# Associative Mapping

- › Direct mapping has high performance only when the cache size is large. (Many blocks may be mapped to the same cache line)
- › In associative mapping, the mapping of the main memory block can be done with any of the cache block.
- › The tag field of the cache line includes the address of the block

# Set-Associative Mapping

- › It is a compromise between direct and associative mappings that exhibits the strength and reduces the disadvantages.
- › Each line includes multiple blocks (hence set-associative)
- › A block is mapped to a specific set. Inside a set, it can be mapped to any line.
- › The most commonly set-associative cache is two-way set associative cache.
- › Set associative mapping is more expensive but has better performance than other mappings

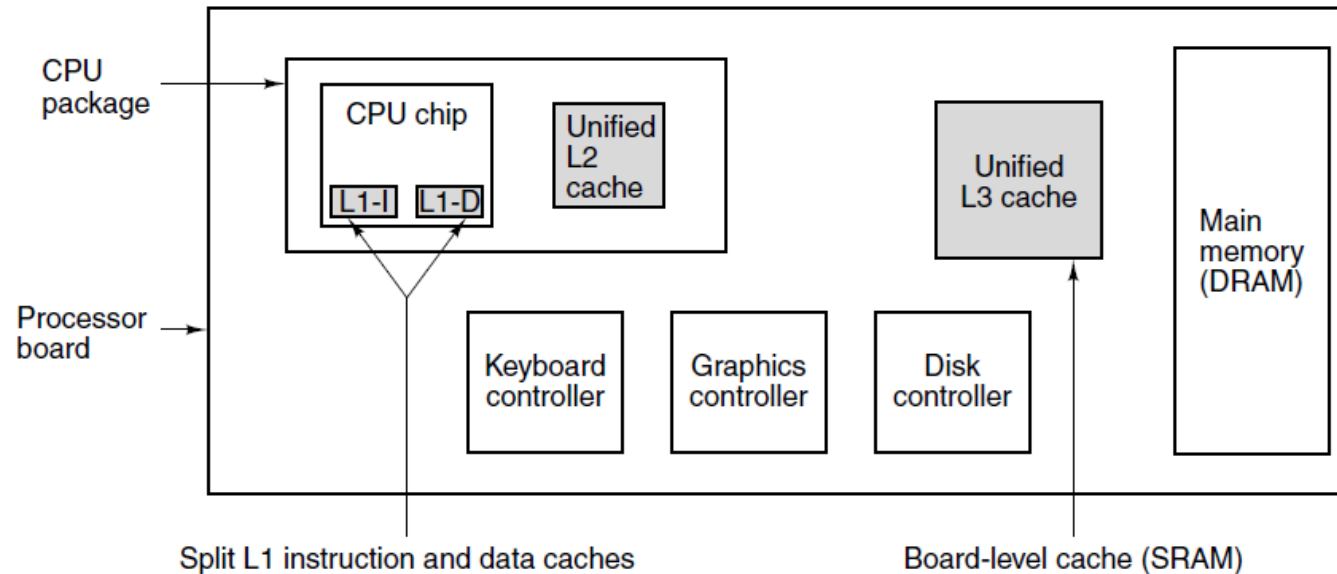
# Example

- › The original Pentium 4 processor had a four-way set associative L1 data cache of 8 KiB in size, with 64-byte cache blocks. Hence, there are  $8 \text{ KiB} / 64 = 128$  cache blocks.
- › Pentium 4 processor also had an eight-way set associative L2 integrated cache 256 KiB in size, with 128-byte cache blocks.

# Cache design

- In practice: fixed-sized memory blocks (64, 32, 16 bytes) referred to as **cache lines**
- Design considerations:
  1. **Size**: bigger = faster, but also more expensive
  2. **Line size**: should we refresh many data or often?
  3. **Administration**: how do we know what is kept inside the cache at any time?
  4. **Content**: unified or split cache for data and instructions?
  5. **Topology**: what is the optimal number of caches? How are they organized?
  6. and many more concerning replacement, power, ...


# Cache memory – hierarchy




- Every cache level from L1 to L3 is bigger but also slower
- Locality principle: **spatial** (A will be followed by A+1) & **temporal** (A will be accessed again in the future)
- Typical cache line sizes of 4 to 64 Bytes



# Intel i7 Cache Memories



# The memory hierarchy



Access time	Storage capacity	US\$ per Bit
Less than nanosecs	Bytes	Dollars per MB
Seconds	(Unlimited)	Pennies per Gb

# (memory) Latencies in perspective

1 CPU cycle	0.3 ns	1 s
Level 1 cache access	0.9 ns	3 s
Level 2 cache access	2.8 ns	9 s
Level 3 cache access	12.9 ns	43 s
Main memory access	120 ns	6 min
Solid-state disk I/O	50-150 $\mu$ s	2-6 days
Rotational disk I/O	1-10 ms	1-12 months
Internet: SF to NYC	40 ms	4 years
Internet: SF to UK	81 ms	8 years
Internet: SF to Australia	183 ms	19 years
OS virtualization reboot	4 s	423 years
SCSI command time-out	30 s	3000 years
Hardware virtualization reboot	40 s	4000 years
Physical system reboot	5 m	32 millenia

<https://twitter.com/srigi/status/917998817051541504?lang=en>



# (longer term note) importance of locality

- Which code has a better locality?

```
int sum_array_rows(int a[M][N])
{
    int i, j, sum = 0;

    for (i = 0; i < M; i++)
        for (j = 0; j < N; j++)
            sum += a[i][j];


    return sum;
}
```

```
int sum_array_cols(int a[M][N])
{
    int i, j, sum = 0;

    for (j = 0; j < N; j++)
        for (i = 0; i < M; i++)
            sum += a[i][j];

    return sum;
}
```

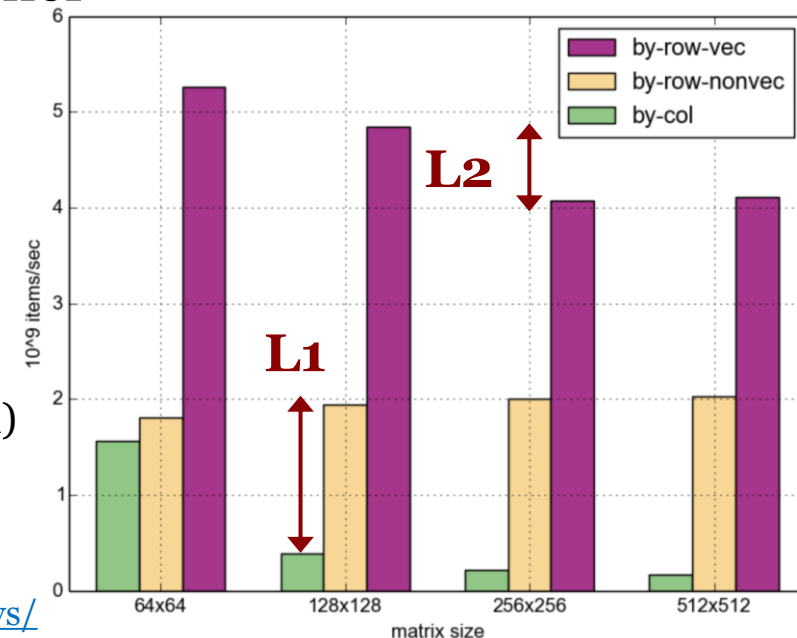
- Being able to look at code and get a qualitative sense of its locality is a key skill for a professional programmer



*Remember to access  
data in order!*

By-row is faster than by-column (6-8x, depends on the size)  
64x64 matrix fits in 32-KB L1 cache (small difference by-row/by-col)  
Larger matrix sizes make by-column go to L2 more frequently  
Vectorized version beats the non-vectorized code by 2-3x  
256x256 and up saturate (L2 (256KB) misses go to main memory)

<https://eli.thegreenplace.net/2015/memory-layout-of-multi-dimensional-arrays/>





university of  
 groningen

# Questions?